Load by Design: Adaptive Heuristics for LLM Checkpoint Loading

Botir Khaltaev
Independent, United Kingdom
btrghstk@gmail.com

Abstract

Inference on large language models is often gated by loading weights from storage into GPU memory. In serverless and multi-tenant settings, idle models are evicted and must be reloaded on demand, so cold-start latency can hinge on how checkpoints are read. Peak NVMe bandwidth is only an upper bound: request structure, checkpoint layout, and software above the block layer often dominate realized throughput.

We present *TensorStore*, an open checkpoint-loading framework and evaluation harness for Linux that supports SafeTensors and a partitioned ServerlessLLM-style layout, with eager backends spanning synchronous POSIX I/O, Tokio-based async loading, and io_uring. Python bindings and a custom vLLM integration extend the study beyond Rust-only microbenchmarks. Experiments use six public checkpoints and both layouts.

Findings. For **contiguous SafeTensors**, explicit **synchronous** loading wins at the small end of our ladder, while explicit **io_uring** wins at the largest multi-shard scales where queue depth matters. No single fixed backend tracks that crossover, so an **adaptive** `default` policy (selector plus planning) is justified; it is not equivalent to any one explicit column. For **partitioned ServerlessLLM**, synchronous loading is never fastest in our matrix; the relevant trade-off is among async, io_uring, and `default`. Python adds enough overhead to reorder some Rust-level conclusions unless the bindings are tuned. In vLLM, loader choice affects initialisation and time-to-first-token far more than steady-state decode. The operational recommendation is a format-aware adaptive default with explicit overrides, not one universal backend. Appendix A documents harnesses, result paths, and environment fields to record for replication.

Keywords: large language models, checkpoint loading, I/O policy, io_uring, cold start, vLLM, reproducibility

Naming. *TensorStore* denotes this paper's checkpoint-loading system (Rust crate `tensor_store`, Python `tensor_store_py`). It is not affiliated with Google's TensorStore library for N-dimensional array storage.

1 Introduction

Serving a large language model begins long before the first token: weights must reach GPU memory, and the route from disk to tensor buffers is easy to underestimate. A model with seven billion parameters in float16 requires roughly fourteen gigabytes of storage and a similar amount of RAM or VRAM to hold the weights alone. In serverless and multi-tenant inference environments, GPU memory is a shared resource. When a model is not actively serving requests, keeping it resident wastes memory that could serve other tenants. Systems therefore evict idle models and reload them on demand.

When a request arrives for an evicted model, the inference pipeline cannot begin until the checkpoint is fully loaded into memory. For large models this loading step takes time. In a production serving environment, this manifests as cold-start latency: the delay between a request arriving and the model being ready to generate the first token. For latency-sensitive applications this delay is directly visible to users.

Modern storage devices offer substantial raw bandwidth. Despite this, checkpoint loading in practice often achieves a fraction of the available throughput. The gap between peak hardware capability and effective loading performance suggests that the bottleneck lies not only in the storage medium itself, but also in how loading requests are submitted, batched, and integrated into the surrounding software stack.

We investigate the checkpoint loading path as a systems problem: how backend choice affects throughput and latency for real LLM checkpoints, how that effect changes with checkpoint layout and model size, and how much of the backend-level gain survives once Python bindings and a serving engine are included.

What prior work does not settle. Format work centres on safe layouts and interoperability; cold-start systems research often optimises *whether* and *when* loads happen (warm pools, snapshots, placement). Database and storage studies show io_uring wins are workload-dependent. None of these strands, alone, supplies operators with *controlled, end-to-end* evidence for *how* to submit reads for modern LLM checkpoints on Linux when bytes must flow through Rust tensor buffers, Python bindings, and an engine such as vLLM. This paper fills that gap with a unified framework, multi-layer measurements, and an adaptive policy grounded in observed crossovers rather than API labels.

We pursue five goals: (1) build *TensorStore* as a unified framework with multiple I/O backends and formats under one interface; (2) implement and evaluate synchronous POSIX I/O, Tokio async loading, and Linux io_uring under that interface; (3) develop a reproducible benchmarking methodology with true cold-start evaluation (page cache dropped between runs); (4) expose the framework through Python bindings and evaluate it with vLLM as a serving-engine case study; and (5) analyse results and produce an adaptive loading heuristic rather than only a backend ranking table.

Mapping goals to evidence. Goals (1) and the core of (2) are substantiated by the TensorStore design and Rust matrices (Sections 4 and 7). Goal (3) is substantiated by the cache-control procedure and anchor repetitions (Section 6). Goals (4) and (5) are substantiated by Python and vLLM layers and by the discussion of regime behaviour and policy (Sections 7–11).

Rust SafeTensors data support a simple explicit-backend pattern: **sync fits smaller fixtures in Table 3; io_uring fits the largest multi-shard cases**. No one backend wins the whole ladder, so **adaptive heuristics** replace a global default. Section 3 separates that from ServerlessLLM, where sync never leads the measured matrix.

Adaptive heuristics means mapping observable features (format, bytes, shard and partition structure) to a backend choice and planning parameters, rather than shipping one static loader. The optimum shifts with scale; the evaluation also reports medium-scale rows where the selector still trails explicit sync (Section 9), because those rows bound honest claims.

TensorStore parses checkpoint metadata into a normalised byte-range read plan, then executes it through selectable backends (POSIX sync, Tokio async, Linux io_uring, plus an adaptive `default`). The same logical plan runs on every backend so differences reflect submission strategy and planning, not skipped work. Section 4 details the architecture; here we only note that SafeTensors (contiguous shards) and ServerlessLLM-style partitioning (fragmented ranges) stress different parts of that stack, which is why regime behaviour differs by layout. Python bindings

and a vLLM integration (Section 7) test whether Rust-level ordering survives binding and serving overhead, including comparison to vLLM's stock (`native`) loader.

We make four contributions. **(1)** TensorStore: a Rust framework with a shared format layer and pluggable backends (sync, async, io_uring), plus harnesses that keep measurement logic off the hot path. **(2)** Empirical regime maps on public checkpoints: a **SafeTensors crossover** between explicit sync at smaller scales and explicit io_uring at the largest scales, and a **ServerlessLLM** matrix where sync never leads. **(3)** Integration evidence: Python binding costs large enough to matter, and vLLM results that isolate loader effects in initialisation and TTFT from decode. **(4)** An adaptive `default` policy and a clear operational takeaway: select by format and scale, keep explicit overrides, and avoid a single global backend.

Artifacts. The Rust implementation (`tensor_store`), Python bindings (`tensor_store_py`), profiling binaries, shell-driven cold-cache harnesses, and result layouts needed to reproduce ordering are documented in Appendix A. Replication should target regime behaviour, not millisecond-identical runs on different hardware.

The remainder of the paper develops the related work and thesis (Sections 2–3), the TensorStore design (Section 4), methodology and results (Sections 5–7), discussion and limitations (Sections 8–9), brief professional considerations (Section 10), and the conclusion (Section 11).

2 Related Work

This section reviews the relevant literature across four areas. First, checkpoint formats and how their storage layout affects loading behaviour. Second, serverless and multi-tenant inference systems and the cold-start problem. Third, the Linux I/O interfaces available to applications, with particular attention to io_uring. Fourth, serving systems such as vLLM and the role of checkpoint loading in serving pipelines.

2.1 Scheduling research, format design, and submission policy

Three strands of prior work are often read in isolation; together they locate our contribution. **Cold-start scheduling and placement** Fu et al. (2024); Silva and Pereira (2021) treats checkpoint loading as a cost to avoid or hide through warm pools, migration, or snapshots. That work optimises *whether* and *when* a load happens; it does not, by itself, answer how fast bytes move through the stack once a load is required. **Format and access-pattern design** (SafeTensors, GGUF, partitioning) determines what byte ranges exist and whether access is sequential, shard-parallel, or fault-driven; it shifts the right loader strategy but does not replace execution policy on top of a fixed logical workload. **I/O interface research** (including database studies of io_uring Jasny et al. (2025)) shows that "use io_uring" is not universal: benefits track queue depth and workload shape. Checkpoint loading inherits that structure, which motivates *adaptive* backend selection keyed to observable regime features rather than a single API label.

Positioning. This paper is not a replacement for scheduling research or for format standards. It complements them by holding formats fixed and measuring *read submission policy* end to end (Rust, then Python, then vLLM) on public checkpoints. Where prior work asks how to avoid loads, we ask how to execute unavoidable loads efficiently when layout and scale move the optimum between sync, async, and io_uring.

2.2 Checkpoint Formats and Loading Behaviour

LLM weights are stored on disk in a structured format. The format determines how loading software reads tensor data: whether reads are sequential or scattered, whether they require many file operations or few, and whether the layout helps or hinders the storage device in maintaining high queue depth.

SafeTensors is a format introduced by Hugging Face for safe model distribution Face (2024). The file begins with a JSON header describing each tensor, including its name, dtype, shape, and byte offset. All tensor data follows the header contiguously. Once the header is parsed, loading software can read each tensor by seeking to its declared offset and reading the declared number of bytes. The format has two properties that are relevant here. First, it avoids loading the entire file into memory before parsing, unlike Python pickle. Second, it prohibits unsafe deserialisation, which matters when loading untrusted checkpoints. From a systems perspective, however, SafeTensors leaves one important question open: it standardises the layout, but not the loading strategy. The format tells the loader where the bytes are, but not how aggressively it should submit reads to the operating system.

SafeTensors does not, by itself, imply a single optimal I/O API. A tiny model checkpoint might load fastest with straightforward blocking reads and a handful of threads, because the work completes before complex submission machinery amortises. A very large checkpoint might only approach hardware bandwidth if the loader can batch and overlap enough requests to feed NVMe parallelism. The literature on SafeTensors emphasises safety and distribution; this paper complements that by treating the *execution policy* as a variable to be measured rather than assumed.

PyTorch checkpoint files typically use Python pickle serialisation. A model is saved by calling `torch.save()`, which serialises both the model state and the pickle protocol metadata. This has well-known safety problems: a malicious checkpoint can execute arbitrary code on load Face (2024). For large checkpoints, loading the serialized file and reconstructing tensors can require substantial host memory, making it impractical for models that approach the available RAM. SafeTensors was designed partly to address both concerns.

GGUF is a more recent format designed for quantised models and memory-mapped loading. It stores tensors in a layout optimised for direct memory mapping and supports multiple quantisation schemes. GGUF tensors are organised so that the memory layout of the file mirrors the layout the model expects at runtime, reducing or eliminating a copy step. This is a different trade-off from SafeTensors: GGUF prioritises memory-mapping friendliness over simple sequential loading.

Partitioned checkpoint formats distribute tensor data across multiple files. ServerlessLLM uses this approach to support partial model loading: only the files needed for the current request are loaded, rather than the entire checkpoint Fu et al. (2024); Contributors (2024). A central index maps each tensor name to its partition file and byte offset. Loading a complete checkpoint requires opening each partition file and reading from the specified offset. The advantage is that loading can be parallelised across files and across network boundaries if partitions are distributed. The disadvantage is that the access pattern becomes more fragmented, with per-file overhead accumulating. This makes partitioned formats especially interesting from a backend-selection perspective, because fragmented range reads stress the submission path differently from large contiguous shard reads.

When evaluating such layouts, it is easy to import intuitions from contiguous formats incorrectly. A partitioned load may issue many small range reads across different file descriptors; that can change which backend wins even at similar total byte counts. We therefore treat ServerlessLLM as a separate regime in the results and do not assume that the SafeTensors crossover pattern transfers without evidence.

What these formats share is that loading is fundamentally a series of read operations with known

byte ranges. The storage system sees a set of read requests. Whether those requests are processed efficiently depends on how they are submitted.

The existing literature describes format design well but focuses on safety, portability, and memory layout. It does not systematically examine how the I/O submission path affects the effective loading speed of checkpoints that are already in a safe format. We address that gap by keeping the checkpoint abstraction fixed while varying the execution policy used to fetch the underlying bytes.

One limitation of treating formats only as static artefacts is that it encourages "checklist" engineering: pick SafeTensors, ship a loader, move on. That checklist does not ask whether the loader's submission model matches the checkpoint scale, shard structure, or integration surface. The critical point for this work is that *safe layout is necessary but not sufficient* for fast cold loads: the same bytes can be fetched under different latency profiles depending on backend choice and planning.

2.3 Serverless and Multi-Tenant Inference Systems

Serverless inference platforms such as AWS Lambda, Google Cloud Functions, and specialised services run LLM inference on demand. Memory and GPU resources are allocated when a request arrives and freed when the request completes. Between requests, resources may be handed to other tenants. When a model is not resident in memory, the platform must load it from persistent storage before inference can begin. This delay is the cold-start latency.

Cold-start latency in serverless settings has been studied from several angles. One class of solutions addresses it through scheduling and placement: keeping models warm, pre-loading them on a pool of standby instances, or migrating running instances to avoid a load altogether Fu et al. (2024); Contributors (2024). Another class addresses it through model optimisation: quantisation, pruning, and early exit strategies that reduce the amount of data that must be loaded or processed. A third class, which this work falls into, examines whether the loading mechanism itself can be made faster.

Snapshot-based warm-up techniques preload a model into memory before the first request arrives Silva and Pereira (2021). These approaches treat the checkpoint loading step as a fixed cost and work around it rather than reducing it.

The cold-start problem becomes more acute as model sizes grow. A seven-billion-parameter model in float16 requires roughly fourteen gigabytes of storage and a similar amount of RAM or VRAM. Loading this much data from an NVMe SSD takes hundreds of milliseconds to several seconds depending on the device and submission path. In a production serving system, this is a large fraction of the total latency budget, particularly for interactive applications.

vLLM is a high-throughput inference engine that supports paged memory management and continuous batching Kwon et al. (2023); vLLM Contributors (2024). Its design is representative of a broader class of serving systems, including ORCA Yu et al. (2022), where throughput improvements come from scheduling, batching, and runtime memory management rather than from a specialised checkpoint-loading path alone. This makes vLLM a useful integration target here: if backend-level gains matter in practice, they should survive at least partially into a serving engine with its own substantial initialisation overhead.

Most cold-start research focuses on what happens before loading begins or after it ends. Loading speed itself is often treated as a constant or folded into a broader serving metric. We instead treat loading as a first-class systems problem and ask which backend policy best matches a given checkpoint regime.

2.4 Linux I/O Interfaces

Applications running on Linux have several interfaces for reading from storage.

The most common is POSIX synchronous I/O, based on the `read()` system call. A thread calls `read(fd, buf, n)` and blocks until the data is available in the buffer. To issue multiple reads concurrently, an application must use multiple threads. Each thread is blocked while its read is in flight. Thread creation and context switching have costs, so thread pool sizes are limited. This limits the number of outstanding reads that a synchronous application can maintain at once.

Linux AIO was introduced to provide genuine asynchronous I/O without threads. However, Linux AIO has significant practical limitations. It does not support many common Linux filesystems and relies on a compatibility layer for others. It requires manual error handling at a low level. The API is complex and brittle to use correctly Kerrisk and contributors (2024a). Few production applications use it directly for these reasons.

io_uring was introduced in Linux kernel 5.1 and designed to address the limitations of earlier async I/O interfaces Axboe (2019); Kerrisk and contributors (2024b). Instead of system calls for each operation, applications interact with io_uring through two shared-memory rings: a submission queue and a completion queue. Submission queue entries (SQEs) are prepared in memory and submitted with a single `io_uring_enter` system call, or in kernel-polling mode with no system call at all. Completion queue entries (CQEs) are retrieved by polling the completion queue.

io_uring provides several features relevant to checkpoint loading. Registered buffers are pre-allocated and registered with the kernel, eliminating per-request memory allocation overhead. Fixed buffers can also reduce buffer-setup and pinning overhead. Batched submission allows many requests to be prepared and submitted together, reducing per-request syscall overhead. Polled completion allows an application to busy-wait on completions without system call overhead when latency is critical. The low-level Rust ecosystem mirrors this design space: Tokio provides a mature async runtime Contributors (2025a,b), while the `io-uring` crate Project and io-uring Contributors (2025) exposes the lower-level ring interface needed to experiment with more explicit submission strategies.

Recent work has examined when io_uring helps in database systems. Jasny et al. studied io_uring integration into a buffer manager and a network-shuffling workload, finding that io_uring benefits depend heavily on workload characteristics and queue depth Jasny et al. (2025). Their results show that io_uring helps most when there are many small to medium-sized requests and the device queue depth is high. For bulk sequential reads, the advantage is smaller because the syscall overhead is already low relative to the data transfer time. They note that naive replacement of synchronous I/O with io_uring does not automatically yield improvements; the workload must exercise the features that differentiate io_uring. This observation is one of the key literature anchors for this paper: backend family alone is not enough, and execution strategy must be matched to workload shape.

This observation is directly relevant to checkpoint loading. A contiguous SafeTensors file with large sequential reads may not benefit as much from io_uring as a fragmented partitioned layout with many small reads. Section 7 examines this question directly, and the final results confirm that backend selection must be heuristic and workload-aware rather than fixed.

The same literature also motivates separating **small** and **large** checkpoint regimes even when both use the same API surface. A small checkpoint may complete in a regime where syscall and setup costs matter in the denominator; a large checkpoint may spend long enough in I/O that keeping the device fed dominates. That distinction is not philosophical; it is exactly the SafeTensors ladder behaviour in our results, expressed without pretending that a single mechanism must win at every scale.

The io_uring interface itself has limitations relevant to checkpoint loading. Registered buffers

work best when all buffers are the same size, but checkpoint tensors vary. Managing variable-sized buffers with fixed registration requires splitting tensors into fixed-size chunks, which adds complexity. The kernel-polling mode consumes CPU even when the device is idle. The interrupt-driven mode has lower CPU overhead but higher latency per completion. There is no single optimal mode for all checkpoint layouts.

NVMe hardware and queueing theory both emphasise the same constraint for application writers: peak media bandwidth matters only if the host keeps submission queues deep enough for the device to work Organization (2024); Hennessy and Patterson (2012). That supports viewing checkpoint loading as a *submission* problem, not only a bandwidth problem.

Memory-mapped access, the page cache, and O_DIRECT are alternative mechanisms with different trade-offs (fault cost, readahead sensitivity, alignment rules). This paper's eager-loading experiments focus on explicit read paths with controlled cold-cache drops (Section 6). mmap-heavy formats such as GGUF are treated in the synthesis subsection as a complementary optimisation target, not a substitute claim set for eager SafeTensors loads.

2.5 Submission policy versus peak bandwidth

Much systems writing implicitly treats "faster storage" as the lever for faster loads. NVMe devices do offer high sequential bandwidth, but application-level throughput also depends on how many operations are in flight and how costly each submission is. A loader that issues reads one at a time from a single thread may leave the device underfed even when the SSD is capable of far more. Conversely, a loader that adds coordination overhead without increasing useful parallelism can be slower than a simpler synchronous design on small files where setup dominates.

This tension is not unique to machine learning. Database research on io_uring emphasises that benefits are workload-dependent: bulk sequential reads already amortise syscall cost, while fragmented or many-small-request workloads are more likely to benefit from batched asynchronous submission Jasny et al. (2025). Checkpoint loading inherits the same structure. Contiguous shard reads resemble bulk sequential access once each shard is large enough; highly partitioned layouts resemble many-range workloads where grouping and queue depth matter more. The open question for LLM serving is therefore not "which API is theoretically fastest," but *which submission policy matches the checkpoint regime at hand*, including how model scale changes the regime.

2.6 Synthesis: GGUF, mmap, and database parallels

Three strands of prior work help position this paper without duplicating its empirical chapters. **GGUF and mmap-first serving assumptions** prioritise eliminating copies by mapping bytes that already sit in a file layout aligned with runtime tensor views. That strategy can win when the operating system's readahead and fault path match access order, and when the working set fits comfortably in virtual memory. It is a different optimisation target from **eager bulk fetch** of SafeTensors shards, where the question is how fast bytes arrive from NVMe through explicit readers. The two are not interchangeable claims: mmap-heavy results do not invalidate synchronous or io_uring eager loads on the same GPU, because the dominant costs sit in different subsystems (fault handling versus submission depth). We hold the eager-loading question fixed and therefore complements mmap-oriented format discussions rather than contradicting them.

Database io_uring studies supply methodological discipline: replace naive synchronous I/O with rings only after the workload exposes enough concurrency for batching to amortise setup Jasny et al. (2025). The checkpoint analogue is not "always io_uring," but "io_uring when the tensor read graph is wide and deep enough that rings keep the device busy." Small checkpoints finish before deep

queueing pays off; partitioned layouts can reach that regime at moderate byte counts because they multiply file descriptors and ranges. Mapping database findings onto LLM checkpoints is therefore qualitative (queue depth, request size, coordination overhead), not a licence to copy database tuning constants.

Cold-start systems research often treats loading as a cost to schedule around (snapshots, placement, keeping replicas warm) Fu et al. (2024); Silva and Pereira (2021). That literature explains why loads happen; it rarely benchmarks how bytes move through the block layer. Where scheduling hides load latency, optimising loaders matters less; where models must spin up from cold storage on demand, loader policy becomes first-class. Our work is orthogonal to placement work: it targets faster mechanics when loading stays on the critical path.

Taken together, these strands justify an **adaptive** stance. Format papers recommend safe layouts; database papers recommend workload-matched submission; serving papers recommend cluster-level tricks. None of them, on their own, tells an operator which single I/O backend to hard-code for an arbitrary checkpoint on an arbitrary integration surface. The gap is operational: a policy that maps measurable workload features to backend choice under explicit integration constraints, which is what Section 6 evaluates.

2.7 Positioning of this work

Prior work supplies components (safe formats, scheduling around loads, database-style io_uring lessons, serving engines). It rarely answers the operational question *which read submission policy to deploy* for a given public checkpoint when the stack includes Rust, Python, and vLLM-style integration. This paper supplies that missing operational map: explicit crossovers on a contiguous ladder, a different ranking under partitioning, and evidence about where integration overhead swallows or preserves Rust-level ordering.

The contribution is deliberately phrased as a **heuristic** backed by measurements, not as a closed-form optimum. We do **not** introduce a new weight format, prove hub security properties, or optimise distributed training checkpoints. The scope is load mechanics and policy under the integration constraints we measure; Section 9 states external validity bounds.

3 Narrative: Scale, Backend Choice, and Adaptive Loading

This section states claims separately for SafeTensors and ServerlessLLM. **Scale shifts the best backend**; **layout can overturn a one-line rule**; **adaptive selection follows from the tables**, not from API names alone.

3.1 SafeTensors (contiguous)

Among *explicit* backends in Table 3, **synchronous** loading wins at the small end of the ladder: blocking reads with parallel shard loading avoid async runtime overhead, and total bytes stay small enough that submission cost still matters. From Qwen3-8B upward, `io_uring` is the fastest explicit column: larger multi-shard loads need sustained queue depth. That crossover makes a **fixed** global backend (always sync or always io_uring) wrong in principle.

The production `default` policy maps workload signals (format, bytes, shards) onto backends and adds planning (chunking, coalescing). It is *not* a relabel of one explicit column, so it can beat or miss explicit sync/io_uring on individual rows.

3.2 ServerlessLLM (partitioned)

Table 4 is a different regime: **synchronous loading never wins**. The trade-off is between grouped **async** work at smaller and mid-sized range-heavy loads and **io_uring** (including via `default` at larger sizes). Do not import the SafeTensors "sync for small" slogan here without the table.

3.3 Why adaptation, and a minimal mechanistic picture

Crossover location depends on shards, grouping, and higher layers (Python, vLLM), so a single static policy cannot track the optimum.

Small SafeTensors loads finish before async/io_uring setup fully amortises; large loads need enough in-flight I/O to feed NVMe parallelism. The evaluated `io_uring` backend matches that regime by sustaining submission depth. This is consistent with the tables, not a formal queueing proof.

3.4 Reading `default` against explicit columns

"Sync for small SafeTensors" and "io_uring for large SafeTensors" refer to the **explicit** sync and io_uring columns as anchors. The `default` column shows the integrated selector with planning. Credit belongs to integration as much as to API names.

The "sync when small, io_uring when large" line is anchored to the **SafeTensors Rust cold-cache matrix** in this harness, not to every filesystem or host. For ServerlessLLM, claims are limited to what Table 4 shows on the tested ladder.

4 TensorStore Design

This section describes the design and architecture of *TensorStore*, the checkpoint loading framework at the centre of this work.

4.1 Design Goals

TensorStore was designed around four primary requirements. First, the framework must support multiple I/O backends and checkpoint formats under a consistent interface, so that backend decisions can be studied without rewriting the rest of the loading stack. Second, the framework must keep the loader logic separate from the measurement harness, so that instrumentation does not distort the measured path. Third, the design must make its concurrency model explicit, because queue depth, worker fanout, and batching policy are central to the performance question. Fourth, the framework must remain accessible from Python and from a real serving engine, so that backend-level findings can be tested beyond Rust-only microbenchmarks.

These requirements led to a layered architecture that separates format parsing, I/O execution, policy selection, and evaluation.

4.2 System Architecture

TensorStore is structured into three layers: a format layer, a backend layer, and an evaluation layer.

The format layer parses checkpoint metadata and produces a normalised loading workload. The key abstraction is a request-local view of the checkpoint: the format layer determines which file ranges must be read, but it does not commit to how those reads are scheduled. This separation

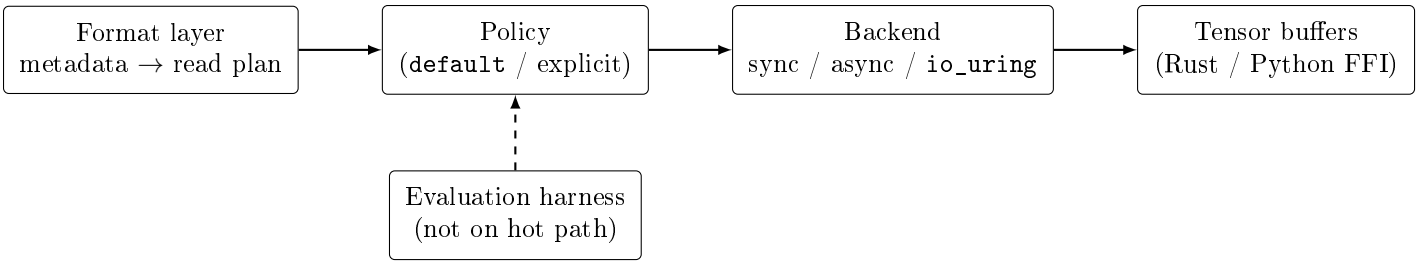

Figure 1: TensorStore architecture: the format layer produces a normalised read plan; policy selects a backend and planning parameters; backends execute I/O; evaluation harnesses attach at the boundary for measurement only.

allows the same logical workload to be executed by different backends while preserving the semantics visible to the caller.

The backend layer implements I/O execution. Each backend receives the same logical workload and executes it according to its own submission model. Backends may perform backend-local planning such as chunking large files or coalescing nearby ranges, but they must preserve request-local output semantics. That matches how real loaders behave: the same logical byte ranges can be fetched with different scheduling while preserving tensor outputs, and comparing backends fairly requires allowing each implementation its natural planning.

The evaluation layer contains the Rust profiling harness, the Python benchmarks, and the vLLM integration benchmarks. These are kept outside the core loading path so that the production implementation and the measurement logic remain distinct.

Figure 1 summarises the layering; the following subsections specify the backend interface, adaptive policy, and bindings.

4.3 Common Backend Interface

All eager-loading backends implement the same public surface:

- `load(path)` for a full file,
- `load_batch(paths)` for multiple whole-file loads,
- `load_range(path, offset, len)` for one byte range, and
- `load_range_batch(requests)` for many ranges.

This interface is deliberately narrow. It is expressive enough for both `SafeTensors` and `ServerlessLLM`, but small enough that backend differences remain interpretable.

4.4 Adaptive Loading Policy

The central design decision is that backend selection is handled as an adaptive heuristic rather than as a fixed policy. Evidence shows that no single backend is best across all checkpoint sizes and layouts.

The final `default` policy uses workload-aware signals and includes backend-internal planning such as chunking large files, coalescing nearby ranges, and workload-driven sizing decisions. This internal planning is part of what makes `default` effective in practice, but it means that `default` is

not a pure selector among the explicit benchmark paths. Instead, `default` represents a higher-level policy that combines adaptive backend selection with execution-time planning decisions.

For `SafeTensors`, the key factors are shard count and total checkpoint size. The adaptive policy is designed around the measured crossover: **keep synchronous execution for smaller eager loads**, and **move to the `io_uring` path as scale increases**. For `ServerlessLLM`, the key factors are total bytes, partition count, and the degree to which the workload is range-heavy; the selector does *not* assume that synchronous I/O is competitive on partitioned range workloads (the Rust matrix shows it is not). Smaller and mid-sized range-heavy workloads may use the Tokio async path; larger partitioned checkpoints use the adaptive `default` that selects `io_uring`. The selector therefore remains format-aware rather than purely size-aware.

The selector is **feature-driven** (bytes, shard and partition structure from the format layer), not a single hard-coded byte threshold. Exact internal cut-offs are omitted here to avoid false precision; regimes matter more than constants. Figure 2 summarises the control flow at a high level; threshold constants live in the open-source implementation alongside the harness revision recorded for the results.

High-level `default` load path (conceptual; names match the codebase structure)

```
features ← extract_features(checkpoint)     // bytes, layout family, shard/partition counts
backend ← heuristic_select(features)     // sync | async | io_uring
plan ← plan_reads(checkpoint, backend)     // chunking, depth, coalescing
return execute(plan)
```

Figure 2: Adaptive policy sketch. SafeTensors and ServerlessLLM use different branches inside `heuristic_select`; `default` may combine selection with extra planning not used by explicit benchmark columns.

Table 1: Major design decisions and rationale.

Decision	Alternatives considered	Rationale
Adaptive `default`	Fixed global backend	Evidence shows crossovers with scale and layout
`io_uring` implementation	Minimal parallelism / naive batching	Worker-owned rings with chunking and planning sustain queue depth on large multi-shard loads
Cold-cache control	In-process eviction hints	External `sync`/`drop_caches` separates harness from loader
vLLM integration	Async loader bridge	vLLM's loader API at the Python boundary is synchronous; benchmarks follow that interface
Python Tokio	New runtime per call	Shared runtime reduced binding overhead

4.5 Backend Designs

The synchronous backend uses blocking POSIX reads. In the final design it is not a purely serial implementation: large files are split into chunks and loaded with host-side thread parallelism. This keeps the backend simple at the API boundary while still exploiting multiple cores where appropriate.

The design intent is not to make `sync` look trivial, but to give it a fair and practical implementation baseline.

The async backend is built on Tokio. Its design is strongest for grouped and range-heavy workloads rather than for large whole-file eager loads. The final implementation groups work per file, uses dynamic batching, and reduces unnecessary task fanout. This backend remains important because the final results show that explicit async loading is often strong for `ServerlessLLM`-style partitioned workloads even when it is weak for `SafeTensors`.

The `io_uring` backend uses multiple worker-owned rings, lazy ring initialisation, dynamic ring sizing, and workload-driven planning for chunk size and in-flight depth, so that submission can keep pace with large multi-shard and range-heavy workloads.

Table 2: Evaluated `io_uring` backend structure.

Aspect	Design
Ring ownership	Worker-owned rings (lazy initialisation)
Submission	Multiple workers; batched SQEs; depth matched to workload
Planning	Dynamic ring sizing, chunk sizing, and in-flight limits

We do not conclude that `io_uring` is inherently superior. We conclude that `io_uring` pays off when the implementation sustains enough parallelism for the workload; topology and planning matter as much as choosing the API family.

The mmap-backed path is retained as a useful lazy-access mode, especially for Python and vLLM experiments where a caller may not need to materialise every tensor immediately. It is not treated as the main eager-loading solution to the cold-start problem, because mmap solves a different problem: efficient selective access rather than aggressive eager loading.

4.6 Checkpoint Formats

TensorStore supports two checkpoint layouts.

`SafeTensors` stores tensors in a whole-file or multi-shard layout. Once the header is parsed, the loader sees a largely contiguous eager-loading workload.

The `ServerlessLLM` layout distributes tensors across multiple partition files with a central index. This introduces a more fragmented range-oriented workload and therefore stresses a different part of the backend design space.

The purpose of supporting both formats is not just compatibility. It is to expose two different I/O regimes under one framework.

4.7 Python Bindings

TensorStore exposes its functionality through Python bindings implemented with PyO3. The bindings intentionally expose a small public API: a `default` eager path, explicit `sync` and `async` eager paths, and lazy `open_*` mmap-backed handles. On Linux, `io_uring` remains an internal backend behind `default` in the public Python API, while benchmark-only explicit paths were used in the vLLM benchmark layer.

One important binding design decision is that the async Python path runs on a shared Tokio runtime rather than constructing a fresh runtime per call. This was one of the practical optimisations that materially improved Python-side performance.

The vLLM integration is implemented through a custom model-loader path. vLLM is particularly useful as an evaluation target because it separates model initialisation, time-to-first-token,

and steady-state decode. This allows us to test whether backend gains survive contact with a real serving engine.

The key design constraint is that vLLM's loader interface is synchronous. The final benchmark integration therefore compares synchronous and default eager-loading policies directly, and uses benchmark-only explicit loader variants where necessary to expose additional backend behaviour. This is a practical design choice: it reflects the serving interface that actually exists, rather than inventing an async serving path that the engine does not support.

Cold-cache control is implemented in the experiment harness, not inside the loader: the Linux shell runs `sync` and `drop_caches` before benchmark invocations. That keeps measurement ritual separate from loading logic and matches common storage-evaluation practice.

This decision is important for clarity. The loader implementation should express how TensorStore loads checkpoints; the benchmark harness should express how cold and warm conditions are created. Tables 1 and 2 summarise principal design choices; the narrative above explains Rust core choice, shared interface, and binding integration in full.

5 Development methodology

This section summarises how the implementation was developed and validated, complementing the technical design in Section 4. The focus is on requirements, quality practices, and how measurements were tied to reproducible harnesses.

5.1 Lifecycle and iteration

Development interleaved implementation with benchmarking so design choices tracked measured behaviour on real checkpoints. Git history links coherent changes (refactors, backend fixes, tests) to the revisions used for final numbers where relevant.

5.2 Requirements and scope management

Initial requirements were to support SafeTensors and ServerlessLLM-style partitioned layouts, expose multiple backends for comparison, and remain usable from Python for integration-style tests. As the codebase grew, secondary requirements appeared: stable error reporting for conversion pipelines, parity tests between backends, and compatibility constraints from downstream tools (for example, vLLM's synchronous loader interface). Scope was managed by treating the Rust core and bindings as separate layers: core correctness and performance first, then binding ergonomics, then serving-engine integration. The vLLM-facing code paths follow the engine's synchronous loader interface at the Python boundary.

5.3 Testing strategy

Testing relied on automated checks at multiple levels. The Rust crate uses unit and integration tests across readers, writers, backends, and converters. The Python bindings use pytest with optional framework extras so that environments without Torch or TensorFlow can still run core tests. Benchmarks and profiling binaries are kept separate from the library surface so that measurement code does not complicate production paths.

Regression risk was controlled by running the full Rust test suite after structural changes and by using continuous integration to catch failures on push. The exact CI configuration may evolve, but the intent is stable: build, test, and lint should remain green on the main development line.

5.4 Quality practices

The codebase adopted consistent error types instead of ad hoc strings for failure modes that users must diagnose (missing tensors, partition size mismatches, and similar). Public APIs document error conditions where non-obvious. Unsafe code used for zero-copy I/O includes explicit safety comments. These practices support long-term maintainability: future contributors can see invariants and failure modes without relying on tribal knowledge.

5.5 Software artefacts

The primary artefacts are a reusable library and bindings with documented loading paths, plus evaluation harnesses used to produce the results in Section 7. Documentation and examples were updated as workflows changed (for example, defaulting conversion steps in helper scripts so that demos fail less often for new users). This aligns engineering effort with the goal of reproducible evaluation rather than one-off scripts.

5.6 Testing and regression discipline

Automated tests served two roles: correctness (tensor bytes match expectations; error paths behave) and guardrails during refactors (backend parity on shared workloads). Benchmarks were treated as non-deterministic measurements: they validate directionally, but headline numbers in this paper come from controlled harness runs with explicit cache-drop procedures rather than from a single CI timing assertion. This separation keeps CI fast and reliable while still allowing honest performance work.

5.7 Maintainability implications of adaptive policy

An adaptive loader is more complex than a single-backend loader. That complexity must be justified by measurable benefit, which is why we tie policy changes to observed crossovers rather than to aesthetics. For future maintainers, the useful invariant is: **when adding a new format or a new backend, update the selector's features and re-run the Rust matrices before claiming parity**. Otherwise the codebase risks drifting into a configuration-heavy system whose behaviour is hard to explain, which would undermine the very clarity that motivates adaptive heuristics in the first place.

Public APIs stay small, explicit backend labels stay available for debugging, and harnesses stay out of library code paths so an adaptive loader stays auditable as maintainers change.

6 Experimental Setup

This section describes the experimental methodology used to evaluate the checkpoint loading backends.

6.1 Constructs at each measurement layer

Rust cold load time isolates format processing and I/O through the selected backend to tensors in the Rust process; it excludes Python materialisation and vLLM CUDA setup. **Python load time** includes FFI and tensor construction and answers what a binding client observes. **vLLM metrics** (`load_only`, `TTFT`, decode) include serving-engine initialisation and are the right construct

for deployment impact, not for pure disk attribution. Claims must match the layer that measured them.

6.2 Research Questions

The evaluation is organised around five research questions:

- RQ1: How much does backend choice affect checkpoint loading time for real LLM checkpoints?
- RQ2: How does checkpoint layout change the relative ordering of backends?
- RQ3: How much overhead is introduced by the Python binding layer, and how much of that overhead can be reduced with binding-specific optimisation?
- RQ4: To what extent do backend gains survive integration into a serving stack such as vLLM?
- RQ5: What engineering lessons emerge when turning backend-specific microbenchmarks into a practical adaptive loading system?

6.3 Experimental Platform

All experiments were run on a dedicated Linux server with an NVIDIA H100 GPU and approximately 180 GB of host RAM, used as an isolated experiment host. The CPU is a contemporary multi-core x86_64 processor (record the exact model next to published timings if your submission requires it).

The software stack for final runs was Ubuntu 24.04, Python 3.12, and Rust 1.94 (`rustc 1.94.0`, stable). Kernel version was recorded beside exports (`uname -r`; 6.x on the measurement host). The optional vLLM dependency is pinned in `bindings/python/uv.lock` (resolved `vllm` 0.11.0); serving exports should still record `pip show vllm` or `uv pip show vllm` next to tables. Appendix A lists other environment fields to capture.

Storage is local NVMe behind a standard Linux journaling filesystem with default metadata behaviour unless noted. Cold-cache runs for the Rust matrix use buffered I/O and external cache drops (procedure below), not O_DIRECT, unless a specific backend documents otherwise.

GPU-backed vLLM runs are constrained by vLLM memory budgeting (KV cache, CUDA graphs, buffers), not only raw VRAM. I/O experiments are Linux-only because the io_uring backend is Linux-specific.

6.4 External baselines

The Rust matrices compare TensorStore's own backends on *identical* logical read plans; they isolate submission policy rather than pitting TensorStore against unrelated third-party loaders.

The vLLM evaluation includes a production-style baseline: the `native` loader uses vLLM's stock checkpoint path for the same model configuration and cache mode. Loaders prefixed `ts_` are TensorStore-backed integrations (`ts_safetensors_*`, `ts_serverlessllm_*`) through the same harness so that differences isolate the weight-loading path. Table 6 summarises cold `load_only` times where `native` is available.

At the Python layer, benchmarks target the TensorStore public API (explicit backends and `default`). A conventional Hugging Face workflow often uses framework-native tensor construction on top of `safetensors`; that stack differs from our Rust-first path, so we report TensorStore binding stages (Table 5) rather than claiming a single universal Python baseline number.

6.5 Backends Under Test

The system exposes four eager-loading policies at the Rust layer:

- **sync**: blocking POSIX reads with dynamic chunking and parallelism.
- **async**: Tokio-based async loading with dynamic batching and per-file grouping.
- **io_uring**: a low-level Linux `io-uring` backend with worker-owned rings, dynamic ring sizing, and dynamic chunk planning (Section 4).
- **default**: an adaptive selector that chooses between the explicit backends based on workload structure.

Rust-layer timings for `io_uring` refer to that implementation (Section 4), not to the syscall family in the abstract.

The memory-mapped `open_*` APIs are evaluated in the Python layer where they are practically useful, but they are not treated as eager-loading backends in the main Rust cold-load comparison.

Two checkpoint layouts were evaluated:

- **SafeTensors**: whole-file or multi-shard eager loading.
- **ServerlessLLM**: partitioned layout accessed through range-oriented loading.

The model ladder consists of six open-access checkpoints:

1. `Qwen/Qwen3-0.6B`
2. `HuggingFaceTB/SmolLM3-3B`
3. `Qwen/Qwen3-4B`
4. `Qwen/Qwen3-8B`
5. `Qwen/Qwen3-14B`
6. `Qwen/Qwen3-32B`

This ladder was chosen to cover small, medium, and large checkpoints while also including one cross-family model (SmolLM3-3B) to reduce the risk of overfitting conclusions to a single vendor family.

The ladder is also what makes the SafeTensors narrative empirically meaningful. A single-point benchmark at one model size cannot justify adaptive policy; it can only justify a constant backend choice. By spanning orders of magnitude in total bytes and shard structure, the experiments expose where synchronous loading stops being the best explicit option and where io_uring becomes the best explicit option in Table 3. That progression is the evidential core behind Section 3.

Cold-cache behaviour is the primary condition for the Rust and Python eager-loading measurements. Cold-cache runs use a manual Linux-side procedure before each benchmark invocation (no in-loader eviction):

1. run `sync`
2. run `echo 3 > /proc/sys/vm/drop_caches`

3. invoke the benchmark binary or Python benchmark

Warm-cache conditions skip the cache-drop step and reuse the existing page-cache state.

Rust-layer measurements were generated with a standalone profiling harness. Two categories of Rust experiments were used. First, a five-repetition cold-cache anchor set was collected for Qwen3-0.6B, Qwen3-8B, and Qwen3-32B for both SafeTensors and ServerlessLLM, across sync, async, io_uring, and default. These repeated measurements were used to identify stable workload regimes and to set default-selector thresholds. Second, a full cold-cache matrix was run across the six-model ladder, both formats, and the four eager backend policies.

Python benchmarks were run against the public binding surface rather than against every internal Rust backend directly. This reflects the actual intended Python API: adaptive default policy, explicit synchronous eager load, explicit Tokio eager load, and mmap-backed lazy access path. The main Python benchmark model was Qwen/Qwen3-8B.

The Python layer was benchmarked across stages that separate pure backend effects from binding-layer tuning and selector behaviour (Table 5).

vLLM benchmarks were run through a custom benchmark loader integration. Unlike the Python eager-loading benchmarks, the vLLM evaluation measures the loading backends inside a wider serving-engine context.

Three benchmark kinds were collected:

- **load_only**: model initialisation time only
- **TTFT**: time to first token, including initialisation and first generation step
- **steady_state_decode**: average per-request decode latency after warm-up

Both warm and cold cache conditions were collected. The full vLLM matrix covers all six fixtures and all available benchmark loaders.

Three rules were used to keep the evaluation interpretable.

1. Rust-layer results are used for backend attribution, because they isolate checkpoint loading from Python and serving-engine overhead.
2. Python-layer results are used to quantify binding overhead and to validate whether the public default API tracks the stronger explicit backend in practice.
3. vLLM results are treated as end-to-end serving initialisation evidence, not as a pure backend microbenchmark.

Backend gains do not propagate unchanged across layers.

6.6 Operational definitions and internal validity

What "load time" means at each layer. At the Rust layer, the profiling harness measures the elapsed time to read checkpoint bytes from storage through the selected backend and assemble tensors in memory according to the format's semantics. It intentionally isolates that work from Python object construction and from GPU allocation inside a serving engine. At the Python layer, reported times include interpreter overhead, tensor materialisation, and any FFI transitions exposed by the binding API; they answer whether a real client of the library would observe backend differences, not what a pure Rust microbenchmark would print. At the vLLM layer, `load_only`

and `TTFT` include engine initialisation steps beyond tensor bytes on disk: CUDA context setup, memory pools, and runtime configuration. Steady-state decode measures post-warm-up per-request latency and is dominated by compute and scheduling; it is included to show where loader choice stops mattering, not to rank I/O APIs for decode.

Cold cache versus warm cache. A **cold-cache** run is one where the checkpoint bytes are not resident in the Linux page cache at the start of the measurement. The procedure used here is `sync` followed by `echo 3 > /proc/sys/vm/drop_caches`; that sequence drops clean cache pages system-wide before invoking the benchmark. That is stronger than evicting a single file with advisory hints and matches common practice in storage evaluation. It does not guarantee an empty GPU or an idle disk firmware queue; it guarantees that the loader cannot satisfy reads entirely from DRAM backed by a prior read of the same file. A **warm-cache** run skips the drop and reuses pages left by earlier reads on the host, modelling a second load without a full cache flush between runs.

Why eviction is out-of-process. In-process `posix_fadvise` and similar hints are advisory and filesystem-dependent; embedding them in the loader would blur the line between harness and code under test. External cache drops keep the loader binary unchanged: every cold iteration pays the same visible storage path, and the methodology is easy to replicate from a shell script.

Variance and anchor repetitions. NVMe and kernel subsystems exhibit run-to-run jitter. Single cold numbers are therefore *representative* rather than statistically exhaustive. The five-repetition anchor set exists to confirm that observed crossovers are not one-off artefacts for a handful of ladder points (small, large, and mid-scale anchors across formats). The full matrices use one cold iteration per cell for readability. Interpreting a 5% gap as meaningful without repeats is risky; interpreting a consistent ordering across many cells and formats with anchors showing stable bands is much less so.

Comparability across backends. Each backend receives the same logical tensor read plan for a given fixture and format: the same sequence of byte ranges must be satisfied. Differences in wall time come from submission strategy, parallelism, and buffer management, not from silently skipping work. When `default` loses to an explicit backend on a row, that is an empirical fact about the selector and planner under the same harness constraints, not a different problem instance.

Statistical reporting. Full matrices use representative cold-cache times; anchor cells use five repetitions to check variance. Appendix B reports median wall-clock time and the min–max band for selected anchor configurations (SafeTensors and ServerlessLLM). Main tables omit per-cell standard deviations for readability, so small gaps in one cell should be read cautiously, while consistent ordering across many cells carries more weight.

6.7 Reading the matrices and reproducing ordering

Read Rust tables alongside Section 3. SafeTensors compares sync, async, io_uring, and `default` on the same logical workload; `default` shows when integrated planning beats or misses explicit backends. The ladder orders increasing bytes and shard complexity in this harness only. SmolLM3-3B and Qwen3-4B sit in a medium band where `default` still trails explicit sync (Section 9). On ServerlessLLM, treat sync as a baseline, not a target.

Cold-cache runs use `sync` and `drop_caches` so numbers reflect storage, not warm DRAM. vLLM warm-cache cases are labelled separately.

Replication should not chase identical milliseconds across machines; it should reproduce **ordering**: SafeTensors sync→io_uring crossover among explicit backends, and ServerlessLLM with sync off the front of Table 4. Ship revision hashes, model IDs, partition conversion settings, and cold-cache commands. Improvements here are **conditional**, which is why adaptation exists.

7 Results

The evaluation is organised in three layers: Rust (backend attribution), Python (binding cost), and vLLM (serving stack).

The Rust **SafeTensors** ladder shows explicit **sync** at the small end and explicit **io_uring** from Qwen3-8B upward (Table 3), which motivates adaptive policy. The **ServerlessLLM** matrix never has sync fastest (Table 4); the comparison there is async versus io_uring and `default`, not sync versus io_uring.

7.1 Rust-Layer Backend Results

The Rust harness is the authoritative backend-attribution layer because it measures checkpoint loading directly, without Python tensor conversion or serving-engine initialisation costs. Unless stated otherwise, each matrix cell is one representative cold-cache run; Appendix B reports five-repetition medians and min–max bands for selected anchors.

7.1.1 SafeTensors

Table 3 summarises the explicit backend comparison for SafeTensors. Two regimes are immediately visible.

Table 3: Rust cold-cache SafeTensors results (ms).

Model	Sync	Async	Default	io_uring
Qwen3-0.6B	555.28	1416.19	333.59	1159.24
SmolLM3-3B	1331.79	5802.05	3722.35	3812.35
Qwen3-4B	1820.85	5039.21	3083.18	2967.42
Qwen3-8B	3880.40	7074.40	3844.35	3172.37
Qwen3-14B	7030.25	13489.00	6138.70	5977.07
Qwen3-32B	15454.45	28763.46	12119.50	12089.05

For smaller eager whole-file loads, synchronous loading remains highly effective. On SmolLM3-3B and Qwen3-4B, sync is the strongest explicit backend. The ranking flips on larger multi-shard checkpoints. On Qwen3-8B, Qwen3-14B, and Qwen3-32B, `io_uring` is the fastest explicit backend. The resulting rule is the clearest high-level heuristic in our evaluation: sync for smaller SafeTensors checkpoints and io_uring for larger ones.

The async backend is consistently the weakest choice for SafeTensors. This matches the workload structure: whole-file shard loading benefits from strong bulk parallelism, not from scheduler-heavy fine-grained async orchestration.

Figure 3 shows the crossover from sync-favoured smaller models to io_uring-favoured larger ones.

`default` improves on a purely cost-based selector overall but misses some medium SafeTensors rows (SmolLM3-3B, Qwen3-4B); Section 9 discusses that gap. For ServerlessLLM, `default` is strong on Qwen3-8B and Qwen3-32B and tracks async/io_uring regimes elsewhere. The `default` column includes planning absent from explicit paths, so it is not a pure selector test.

7.1.2 ServerlessLLM layout

The ServerlessLLM layout produces a different ordering, shown in Table 4.

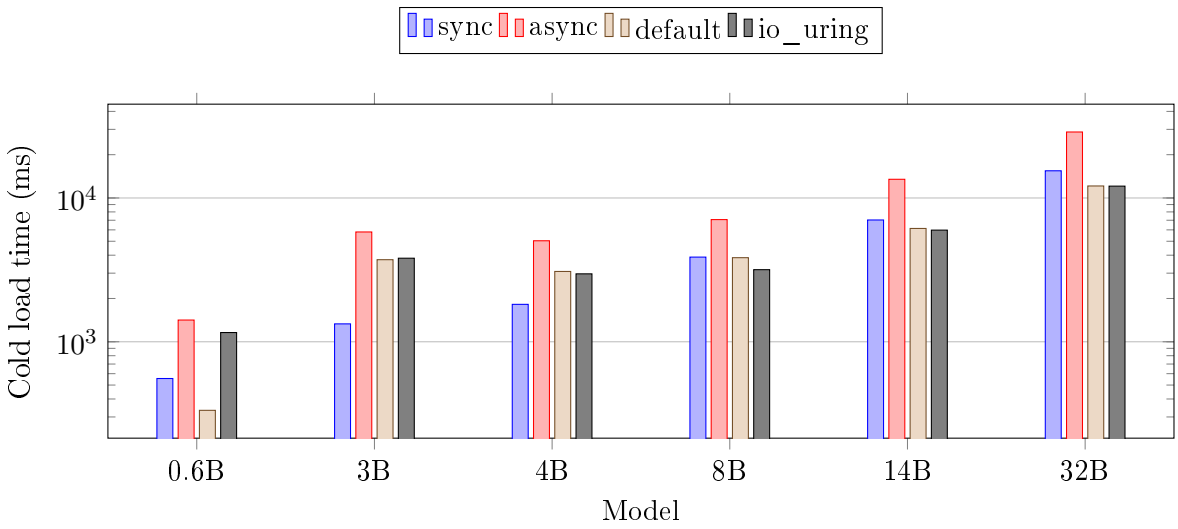

Figure 3: Rust cold-cache SafeTensors crossover behaviour (logarithmic time axis).

Table 4: Rust cold-cache ServerlessLLM results (ms).

Model	Sync	Async	Default	io_uring
Qwen3-0.6B	1645.06	1471.34	1432.25	1387.57
SmolLM3-3B	3941.88	2391.24	3182.28	3091.49
Qwen3-4B	5137.72	3504.25	3551.82	3727.14
Qwen3-8B	11048.40	9463.37	7589.47	9043.51
Qwen3-14B	18661.25	16139.08	13422.94	13257.23
Qwen3-32B	39822.09	29180.89	27659.78	29779.86

Here, sync is never the best backend. For smaller and mid-sized range-heavy workloads, explicit async is often strong because grouped and coalesced range reads fit its execution model well. For larger partitioned checkpoints in this ladder, io_uring becomes competitive or best among explicit columns, and the adaptive `default` policy routes these sizes to the io_uring path (Qwen3-8B and Qwen3-32B are cases where `default` outperforms explicit io_uring because of additional planning). Qwen3-14B is fastest with explicit io_uring; SmolLM3-3B and Qwen3-4B favour async. The partitioned layout therefore changes the ranking and makes a fixed backend recommendation less defensible.

Figure 4 shows the shift from sync toward async and io_uring. The selector moves between async on smaller and mid-sized range-heavy cases and io_uring for larger partitions; `default` follows that hand-off (explicit io_uring can edge out `default` on Qwen3-14B).

7.2 Interpretation

SafeTensors (explicit columns). Sync fits the small end of the ladder; io_uring fits from Qwen3-8B upward; async is weak on this format in the measured matrix. The adaptive `default` column is not a relabel of sync or io_uring: it wins on Qwen3-0.6B with planning and loses to explicit sync on SmolLM3-3B and Qwen3-4B (Section 9).

ServerlessLLM. Sync loses on every row; winners are async, `default`, or io_uring, with `default` often routing toward io_uring at the largest sizes tested.

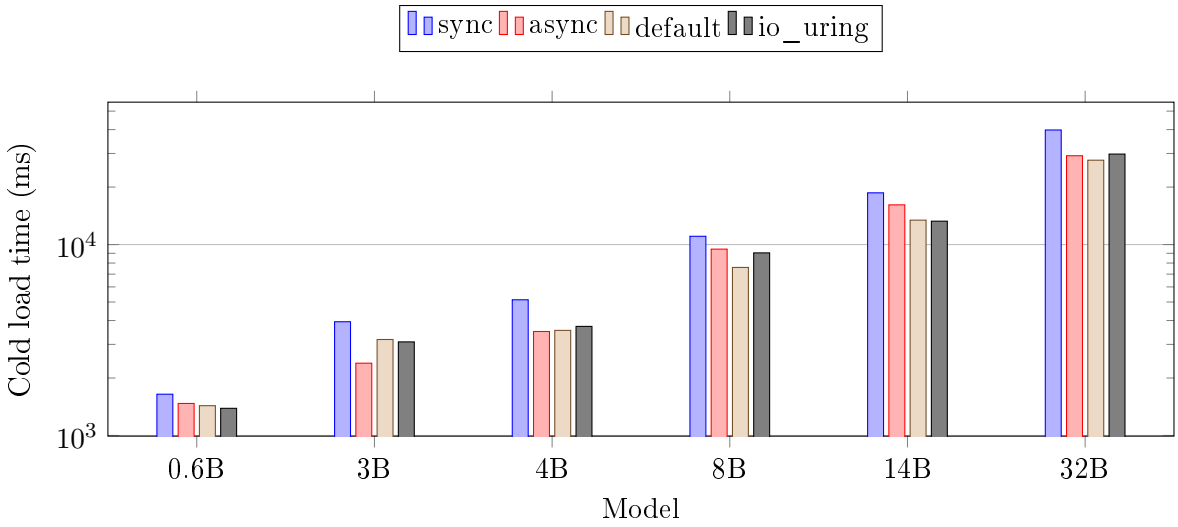

Figure 4: Rust cold-cache ServerlessLLM crossover behaviour (logarithmic time axis).

Layering. Rust tables support backend attribution claims. Python shows tensor materialisation swallowing part of Rust gains unless bindings are tuned. vLLM shows init-time effects under engine memory rules; TTFT adds generation work beyond loading. Integration timings do not replace Rust attribution; they show what operators observe after binding and serving overhead.

7.3 Python Binding Results

Python benchmarks were collected on Qwen/Qwen3-8B to quantify user-facing overhead and to evaluate whether binding optimisation and selector redesign improved practical performance.

The most important Python result is that binding engineering mattered. Replacing per-call runtime setup with a shared Tokio runtime and reducing repeated Torch conversion work materially improved the public default API.

Table 5: Python eager-load times for Qwen3-8B across binding and selector stages (seconds).

Case	Baseline bindings	After binding tuning	After selector tuning
SafeTensors default	15.88	14.55	11.55
ServerlessLLM default	22.45	20.50	18.45

Table 5 shows the sequence. For SafeTensors on Qwen3-8B, the default path moves from 15.88 s to 11.55 s, a 27.3% reduction from the baseline binding column. For ServerlessLLM, the same sequence moves default from 22.45 s to 18.45 s, an overall reduction of 17.8%.

The Python numbers are much larger than the Rust-only numbers because eager loading at the binding layer includes substantial tensor materialisation and Python object construction. The key conclusion is that Python is not a transparent pass-through layer. Backend improvements matter, but they are filtered through the cost of converting checkpoint contents into Torch tensors and Python dictionaries.

7.4 vLLM integration results

The final vLLM matrix completed successfully across all six models, all benchmark loaders, warm and cold cache conditions, and all three benchmark kinds.

Stock vLLM versus TensorStore. The `native` loader follows vLLM's default weight-loading path; `ts_*` loaders use TensorStore-backed integrations (Section 6). On cold `load_only`, TensorStore's default SafeTensors path reduces initialisation time relative to `native` for the fixtures in Table 6, while keeping engine settings fixed. The gap is an end-to-end effect (CUDA, memory pools, and loader policy), not a pure Rust microbenchmark.

Table 6: Cold vLLM `load_only` initialisation: `native` (stock) versus `ts_safetensors_default` (TensorStore). Times in ms; Δ is the relative change in init time (negative means TensorStore is faster).

Model	`native`	`ts_safetensors_default`	Δ
Qwen3-0.6B	23659	20578	−13.0%
Qwen3-8B	56639	35556	−37.2%
Qwen3-32B	169971	88200	−48.1%

For cold load_only, the best loader varies with model size. Table 7 shows the strongest loader in the final cold-load-only configuration.

Table 7: Best cold load_only vLLM loader by model.

Model	Best init time (ms)	Best loader
Qwen3-0.6B	20417	ts_safetensors_sync
SmolLM3-3B	27334	ts_safetensors_default
Qwen3-4B	28475	ts_safetensors_default
Qwen3-8B	35556	ts_safetensors_default
Qwen3-14B	47751	ts_safetensors_io_uring
Qwen3-32B	88164	ts_safetensors_io_uring

For cold TTFT, the ranking is broadly similar to load_only (smaller models favour explicit sync-style loaders; larger models favour default or io_uring), but it does not match row-for-row because TTFT includes the first generation step after initialisation. Table 8 shows the best TTFT loader for each model.

Table 8: Best cold TTFT vLLM loader by model.

Model	Best TTFT (ms)	Best loader
Qwen3-0.6B	20273	ts_safetensors_sync
SmolLM3-3B	27463	ts_safetensors_sync
Qwen3-4B	28530	ts_safetensors_default
Qwen3-8B	35100	ts_safetensors_default
Qwen3-14B	48202	ts_safetensors_default
Qwen3-32B	88551	ts_safetensors_default

Two observations stand out from both tables. First, the best loader for each metric generally

follows size-dependent regimes (smaller models often favour `ts_safetensors_sync`; larger models more often favour `ts_safetensors_default` or `ts_safetensors_io_uring`), but the exact winner can differ between load_only and TTFT for the same fixture (for example SmolLM3-3B) because TTFT adds first-token generation work. Second, the SafeTensors loader family consistently outperforms the ServerlessLLM family for both metrics in this matrix.

Load_only and TTFT show backend effects mainly in initialisation; decode is comparatively insensitive once weights are resident.

At large SafeTensors scales, the evaluated `io_uring` backend sustains utilisation; backend *design* matters as much as API family (Section 4).

1. Rust: structure-dependent ranking; SafeTensors sync small / io_uring large; ServerlessLLM async or io_uring / `default`, almost never sync.
2. Python: binding overhead large; optimisation helped the public default path.
3. vLLM: loader effects concentrate in init and TTFT; best loader varies with model size.

Overall, an adaptive, format-aware default fits the measurements: explicit sync versus io_uring on contiguous SafeTensors by scale; partitioned ServerlessLLM without a sync-first rule.

8 Discussion

Section 3 states the evidence-backed narrative. The following subsections explain why the SafeTensors and ServerlessLLM patterns differ, what changes once Python and vLLM are included, and what practitioners should take away.

No single backend wins across the full ladder and both layouts. The SafeTensors matrix supports explicit io_uring at the top of the ladder and explicit sync on several smaller fixtures, which motivates an adaptive default rather than a single global choice. Python changes ordering relative to Rust; vLLM initialisation dominates while decode is insensitive. Those layers do not erase the Rust crossover; they show where policy must be evaluated.

8.1 Format-specific mechanics

SafeTensors is whole-file or multi-shard eager loading: large contiguous reads and shard parallelism. Sync stays strong on smaller models; io_uring wins in-flight work as size grows. ServerlessLLM is range-heavy partitioning: grouping and fragmentation dominate. Sync is rarely ideal; async is strong on some mid-sized cases; larger partitions favour io_uring and `default`. Backend choice cannot be separated from layout.

Implementation matters, not only API names. For `io_uring`, ring topology and worker ownership determine whether submission keeps CPU and storage busy at scale in the evaluated backend (Section 4).

Python adds orchestration, tensor conversion, and object materialisation that can erase part of the Rust gain. Shared Tokio setup and less redundant Torch work helped Qwen3-8B. Profile bindings separately from the Rust core.

vLLM. Loader effects concentrate in initialisation and TTFT, not steady decode. Reported end-to-end times depend on engine memory budget and `max_model_len` as well as on the loader path; interpret integration results in that full-stack context.

8.2 Policy claims that survive scrutiny

The narrow defensible core is: **on contiguous SafeTensors in this harness, explicit sync wins on smaller fixtures and explicit io_uring on the largest; a fixed backend across the ladder is wrong. ServerlessLLM partitioning breaks any sync-first rule from SafeTensors.** Policy must be format-tagged (contiguous shards versus partitioned ranges), not parameter-count alone. **io_uring is not a guaranteed win from the syscall name alone:** gain depends on whether the implementation matches workload parallelism, not on the API family alone.

Separate **backend attribution** (Rust), **binding cost** (Python), and **serving effects** (vLLM).

8.3 What would change the conclusion

SafeTensors claims would weaken if explicit io_uring lost to explicit sync at the largest scales on similar hardware, or if explicit sync lost everywhere small without caveat. ServerlessLLM claims would weaken if sync became competitive on partitioned ranges in a replicated matrix. Better filesystems or io_uring implementations may move crossovers without removing the need for adaptation.

8.4 External validity and replication

Ordering of backends may generalise more reliably than absolute milliseconds when moving to different NVMe generations, network-attached storage, or filesystems. The structural claim is that policy must track layout and scale; that claim is more portable than any single crossover latency. Replication on other machines should prioritise regime behaviour and ranking, not byte-identical timings.

8.5 Uncertainty, variance, and decision-making

The matrices look definitive because each cell is a single number. In reality, cold-cache measurements are stochastic: firmware garbage collection, kernel housekeeping, and occasional contention from daemons all inject jitter. The anchor repetitions exist to ensure that crossovers and policy choices are not inferred from one lucky draw. When two backends sit close in a cell, treat the ranking as *fragile* unless repeated runs or anchor bands support it; when the gap spans tens of percent across the ladder, treat the regime pattern as *stable* even if milliseconds move.

This matters for operations. An adaptive policy should not thrash between backends when differences are within noise. Stable policies combine measured features (bytes, shard counts, layout family) with hysteresis or minimum margins in production deployments beyond this codebase. The contribution is not a lookup table of millisecond winners; it is evidence that **margins exist** and align with workload structure enough to automate.

Rust, Python, and vLLM differ in **variance sources**: storage and submission (Rust), interpreter and framework jitter (Python), GPU initialisation (vLLM). Noise in one layer does not invalidate another. Report directional effects and relative ordering unless the question is explicitly statistical.

When reporting externally, prefer **bands** from anchors for contested mid-scale rows and **explicit cache-state labels**. Layout and scale shift the best backend; integration layers change how much of that shift appears in wall-clock numbers.

8.6 Practitioner takeaways

1. Prefer an adaptive selector with explicit overrides over one global backend.

2. Optimise the layer that hurts: Rust paths for raw load, bindings for Python wall clock, engine settings for vLLM cold start. When evaluating vLLM, compare against the stock loader as well as custom paths (Table 6).

3. Evaluate format and loader together.

4. When briefing stakeholders, keep formats separate: for SafeTensors, small scale prefers sync and large scale rewards io_uring with adaptation; for ServerlessLLM, partitioned ranges change winners and sync is not a shortcut. Mixing the two briefings contradicts the tables.

9 Limitations

Interpret the claims below with these limitations in mind. They bound what the paper can honestly generalise without additional experiments.

9.1 External validity

Hardware specificity. All final measurements come from one H100-based Linux host with minimal concurrent tenant load. Ordering of backends is more likely to transfer across machines than absolute milliseconds, but that transfer is still an empirical claim. A second site (different CPU or NVMe generation, or cloud block storage) would be the direct test of whether SafeTensors and ServerlessLLM orderings in Section 7 are host-specific.

The storage stack is likewise singular: one local NVMe configuration and one filesystem/mount profile. Different journaling modes, RAID layouts, or network-attached storage would change latency variance and could shift which submission path keeps the device fed. The claim that transfers to **ordering regimes** (small contiguous versus large shard-heavy versus partitioned range-heavy) is more portable than the claim that a particular backend wins by exactly x milliseconds on another site.

9.2 Construct validity

The experiments operationalise "fast loading" differently at different layers, on purpose. Rust timings approximate pure I/O and deserialisation work; Python timings approximate library usability; vLLM timings approximate deployment reality including GPU bring-up. None of these is the same construct as "user-perceived latency for the first chat message in production," which would add authentication, batching, queueing, and network RTT. Section 6 makes those definitions explicit so that claims are not smuggled across constructs.

Loader time is also not the same as "time until the model is profitable to serve." It is one component of cold start. Scheduling, prewarming, and model selection dominate many product stories; we do not measure those larger systems.

9.3 Internal validity and selector claims

Heuristic imperfection. The default selector is not perfect on every Rust matrix row. The medium-size crossover region remains the most difficult area to classify cleanly. For example, on SafeTensors, explicit sync substantially outperforms the adaptive `default` on SmolLM3-3B and Qwen3-4B in the published matrix; stating these rows explicitly avoids overstating selector quality.

Statistical reporting. The full Rust and vLLM matrices report representative cold-cache times after stabilisation; a smaller anchor set used five repetitions for selected models to validate regime

behaviour. The main tables omit per-cell variance for readability, so absolute gaps should be read as indicative rather than as tight confidence intervals.

Serving-stack dependence. vLLM results are sensitive to engine settings such as memory budget and maximum model length. This is an unavoidable property of evaluating inside a real serving stack, but it means that the exact numbers are partly harness-dependent.

Scope. Our evaluation covers two checkpoint layouts and a six-model ladder; it does not exhaust the design space of checkpoint formats or storage hardware.

Scope of claims. The narrative in Section 3 is written to match the tables that were actually measured. It does not prove that synchronous I/O is theoretically optimal for small checkpoints in all possible implementations, nor that io_uring is theoretically optimal for large checkpoints on all Linux versions and all devices. It claims a consistent empirical pattern in a controlled harness and uses that pattern to motivate adaptive selection. Stronger claims would require broader sweeps across hardware, filesystems, and competing loader implementations, which is outside the scope of this work but a reasonable direction for future work.

9.4 What would falsify the adaptive heuristic?

The central heuristic would be undermined if, on similar hardware, explicit synchronous loading remained fastest for the largest SafeTensors checkpoints in the ladder, or if explicit `io_uring` never won at large scale. It would also be undermined if partitioned ServerlessLLM loads showed synchronous loading as fastest across the matrix. The measured evidence supports the opposite ordering patterns in this harness; replication on other storage stacks tests boundary stability, not whether crossovers can exist.

Future work directions suggested by these limitations are practical rather than speculative. Cross-hardware replication would test whether ordering is stable when moving from NVMe Gen4 to Gen5, or from local disks to network-attached storage where latency dominates differently. Filesystem and mount-option sweeps would test whether buffered versus direct I/O changes the SafeTensors crossover point. A broader format study could include GGUF or other mmap-friendly layouts, which may shift the optimum toward lazy access rather than eager loading. None of those directions contradict the claim that **scale and layout change the best policy**; they refine where the boundaries lie.

10 Ethics, safety, and reproducibility

This work evaluates systems behaviour on public model weights; it does not collect or analyse user conversations. The following points matter when deploying the open-source `tensor_store` crate in production.

Provenance and safety. Checkpoints from public hubs should be loaded with the same care as any third-party binary: verify distributor checksums when available, pin revisions, and prefer safe formats (e.g. SafeTensors) over unsafe deserialisation. Kernel-adjacent code (`io_uring`, unsafe Rust) should use pinned dependency versions and regression tests on upgrades.

Privacy and regulation. Benchmarks do not process tenant data; production deployments must still meet access control, encryption, and data-protection requirements where personal data or proprietary weights are involved.

Environmental proportionality. Faster cold loads can reduce wasted idle time but do not address the dominant energy cost of large-model serving (inference and memory traffic). Claims should separate "faster loading" from "lower total footprint" unless measured end-to-end.

Transparency. Performance claims should state cache state, format, and layer (Rust, Python, or vLLM). Negative rows (e.g. where `default` trails explicit sync) are reported in the main matrices and limitations. The name *TensorStore* refers to this work's system (`tensor_store`); it is not affiliated with Google's unrelated TensorStore project for N-dimensional arrays.

Integrity. Upstream engines (vLLM), formats, and community tooling should be cited when results depend on them. Use of generative tools for prose or code should follow institutional policies.

These points align with common professional codes for computing practitioners Society (2023).

11 Conclusion

We return to the research questions posed at the outset.

RQ1: How does backend choice affect checkpoint loading time for real LLM checkpoints?

Backend choice affects loading time substantially, but not uniformly. For **SafeTensors**, the explicit-backend ladder matches the narrative **sync for smaller models, io_uring for larger ones**, which motivates adaptive selection; the `default` column adds planning on top of that handoff. For **ServerlessLLM**, rankings differ: **synchronous loading is never the fastest measured backend**, so the same slogan does not apply; there, explicit async is often strongest at smaller and mid-sized range-heavy loads, and io_uring (including via `default` at larger sizes) is the relevant high-scale tool.

RQ2: How does checkpoint layout affect the relative performance of backends?

Checkpoint layout changes the effective workload. Whole-file shard loading and range-heavy partitioned loading stress different parts of the implementation. SafeTensors rewards strong bulk shard parallelism; ServerlessLLM rewards good request grouping, coalescing, and range scheduling. The evaluation therefore shows that layout and loader must be considered together rather than in isolation.

RQ3: How much overhead is introduced by the Python integration layer?

The Python layer introduces significant overhead through runtime orchestration, tensor conversion, and object materialisation. That overhead is large enough to distort backend-level gains if left untreated. However, binding-specific optimisation reduced the public default path substantially on Qwen3-8B, showing that this overhead can be meaningfully reduced through engineering work at the binding layer.

RQ4: To what extent do backend gains survive integration into vLLM?

Backend gains survive into vLLM most clearly for model initialisation and TTFT. They matter much less for steady-state decode, where GPU compute dominates. The completed vLLM matrix confirms that loader choice still matters in a serving stack, but it also confirms that backend microbenchmarks alone are not enough to predict end-to-end behaviour.

RQ5: What practical engineering lessons emerged?

Backend architecture matters as much as choosing an API family: the evaluated `io_uring` backend relies on worker-owned rings and workload-driven planning to match large-checkpoint parallelism. The appropriate engineering target is therefore not one universal backend, but an adaptive, format-aware default policy backed by explicit override paths.

Claims are tied to operational definitions for Rust, Python, and vLLM timings, cold-cache procedures, and artefact locations (Appendix A). Replication should prioritise ordering and regime behaviour when absolute times differ across hardware.

The adaptive `default` policy is not optimal on every matrix row: on some medium SafeTensors fixtures, explicit synchronous loading still beats `default`, which is why explicit overrides remain

part of the operational design (Section 9).

In sum, **checkpoint loading policy must track scale and layout**. For contiguous SafeTensors, sync fits the small end of our ladder and io_uring fits the largest multi-shard cases; adaptation is required because the crossover is real. For partitioned ServerlessLLM-style loads, sync does not lead in our matrix; async and io_uring (including via `default`) carry the trade-off. The contribution is operational: a format-aware adaptive default with explicit overrides, validated across Rust, Python, and vLLM layers.

A Reproducibility

This appendix summarises how measurements were produced: what to hold constant, where artefacts live, and which commands reproduce the Rust-layer backend ordering. The goal is traceability, not identical millisecond reproduction on other hardware.

A.1 Source bundle and public PDF

For public distribution (for example arXiv), submit the `.tex` sources, figures, bibliography, and `acmart.cls` / `ACM-Reference-Format.bst` if the archive requires a self-contained build. Build with `pdflatex`, `bibtex`, then `pdflatex` twice, or use `latexmk`. When later targeting a conference, replace `nonacm` with the venue's ACM options, add rights metadata, and anonymise as required by the call for papers.

A.2 Repository and revision

The implementation and harnesses live in the same Git repository as this paper. Tabulated timings should correspond to a **specific commit hash**; record `git rev-parse HEAD` next to exported results. Later commits may change timings even when ordering stays similar.

The Rust crate builds a profiling binary `profile` (`cargo build -release -bin profile`). Release mode matters: debug builds are not representative of the optimised paths evaluated in the matrices.

A.3 Rust profiling harness

The Rust layer is invoked as:

```
cargo run --release --bin profile -- <format> <backend> \
  --fixture <FIXTURE_NAME> [--iterations N]
```

where `<format>` is `safetensors` or `serverlessllm`, and `<backend>` is `sync`, `async`, `io-uring`, or `default`. The fixture name refers to a directory under `fixtures/` (for example, `qwen-qwen3-8b`), populated via `scripts/download_models.py` and, for ServerlessLLM, the repository's conversion steps.

Cold-cache runs use `-iterations 1` with the cache-drop procedure in Section 6. The full cold matrix was driven by `scripts/run_full_cold_matrix.sh`, which loops formats, fixtures, and backends, runs `sync` and `echo 3 > /proc/sys/vm/drop_caches` before each invocation, and appends timings to `results/h100/profile/full_cold_matrix.tsv`. Rows that fail parsing remain in the TSV so "crash" is distinguishable from "slow."

Five-repetition anchors for selected models use `scripts/run_anchor_reps.sh` (`results/h100/profile/anchor` to sanity-check variance; main tables use representative single cold-cache rows per cell.

A.4 Python and vLLM harnesses

Python benchmarks call the public binding API under `bindings/python/benchmarks/` (for example `uv run pytest benchmarks/...`), with optional extras when Torch or vLLM are required. The repository provides `scripts/run_benchmarks.sh`, which runs the pytest suites and writes pytest-benchmark JSON (default path under `results/benchmarks/`; set `TENSOR_STORE_BENCH_MODEL` to a Hugging Face model id).

vLLM runs use a custom runner; `scripts/run_vllm_full_matrix.sh` sweeps loaders, fixtures, and kinds (`load_only`, `TTFT`, `steady_state_decode`), with `sync` and cache drops before cold conditions. Outputs live under `results/h100/vllm/`. `scripts/rerun_failed_vllm.sh` retries failed rows after harness fixes.

A.5 Environment pins

Capture at least `uname -r`, Python and Rust versions, vLLM package version, and `nvidia-smi` driver output beside result tables. Kernel and CUDA versions affect `io_uring` behaviour and GPU initialisation. The repository lockfile resolves `vllm` to 0.11.0 when the optional extra is installed through `uv`; still record the active environment's `pip show vllm` line with exports so a reader can match wheels and CUDA builds.

A.6 What replication should check

On different machines, prioritise **ordering** and **regime behaviour**: SafeTensors crossover with scale; ServerlessLLM without sync at the front of the ranking; Python overhead large enough to matter; vLLM initialisation sensitive to loader choice while decode is not. Matching every millisecond is unnecessary for validating the main findings.

A.7 Hub names, fixtures, and cache control

The main text uses public hub names (for example, `Qwen/Qwen3-8B`); scripts use matching filesystem slugs (for example, `qwen-qwen3-8b`). Mixing ad hoc checkpoint paths can change shard counts and invalidate comparisons. Prefer hub checksum verification when available.

Dropping the page cache requires `/proc/sys/vm/drop_caches` (typically privileged). The experiments assumed a dedicated host so cache drops do not disturb unrelated tenants. Run `sync` before `drop_caches` so cold reads exercise storage rather than stale clean buffers.

B Cold-cache anchor statistics

Table 9 summarises five cold-cache repetitions for selected anchors (source: `results/h100/profile/anchor_reps.t`
Median wall-clock time is reported with the min–max band across repetitions. These anchors motivated selector thresholds; the main matrices still use one representative cold run per cell for presentation.

C Artefact layout

Rust results sit under `results/h100/profile/`, Python benchmark JSON under `results/h100/python/` where applicable, and vLLM outputs under `results/h100/vllm/`. That separation preserves the distinction between backend attribution (Rust), binding cost (Python), and serving-stack effects (vLLM).

Table 9: Anchor repetitions: median cold load time (ms) with min–max across five runs.

Fixture	Format	Backend	Median	Min–Max
Qwen3-0.6B	SafeTensors	sync	323.35	322.05–402.93
Qwen3-8B	SafeTensors	sync	3700.68	3647.97–3731.11
Qwen3-8B	SafeTensors	io_uring	3250.29	3160.46–3274.00
Qwen3-32B	SafeTensors	sync	15462.39	15326.77–16143.61
Qwen3-32B	SafeTensors	io_uring	12204.11	11736.60–14388.76
Qwen3-0.6B	ServerlessLLM	default	1356.17	1336.94–1404.77
Qwen3-8B	ServerlessLLM	default	8059.26	7586.28–9143.26

References

Jens Axboe. Efficient i/o with io_uring. `https://kernel.dk/io_uring.pdf`, 2019.

ServerlessLLM Contributors. Serverlessllm. `https://github.com/ServerlessLLM/ServerlessLLM`, 2024. GitHub repository.

Tokio Contributors. Tokio documentation. *Docs.rs*, 2025a. URL `https://docs.rs/tokio/latest/tokio/`. Core async runtime dependency used by the project.

Tokio Contributors. Tokio project repository. `https://github.com/tokio-rs/tokio`, 2025b. Implementation ecosystem reference.

Hugging Face. Safetensors documentation. `https://huggingface.co/docs/safetensors`, 2024. Accessed: April 2026.

Yao Fu, Leyang Xue, Yeqi Huang, Andrei-Octavian Brabete, Dmitrii Ustiugov, Yuvraj Patel, and Luo Mai. Serverlessllm: Low-latency serverless inference for large language models. In *Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 2024)*. USENIX Association, 2024. URL `https://www.usenix.org/system/files/osdi24-fu.pdf`.

John L. Hennessy and David A. Patterson. *Computer Architecture: A Quantitative Approach*. Morgan Kaufmann, 5th edition, 2012.

Matthias Jasny, Muhammad El-Hindi, Tobias Ziegler, Viktor Leis, and Carsten Binnig. High-performance DBMSs with io_uring: When and how to use it, 2025. URL `https://arxiv.org/abs/2512.04859`.

Michael Kerrisk and contributors. *aio(7) Linux Manual Page*, 2024a.

Michael Kerrisk and contributors. *io_uring(7) Linux Manual Page*, 2024b.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP 2023)*. ACM, 2023. doi: 10.1145/3600006.3613165. URL `https://arxiv.org/abs/2309.06180`.

NVM Express Organization. Nvm express specifications. `https://nvmexpress.org/specifications/`, 2024.

Tokio Project and io-uring Contributors. io-uring rust crate documentation. https://docs.rs/io-uring/latest/io_uring/, 2025. Implementation reference for the low-level Linux backend.

Paulo Silva and Thiago Emmanuel Pereira. Performance evaluation of snapshot methods to warm the serverless cold start, 2021. URL https://arxiv.org/abs/2105.13894.

British Computer Society. Code of conduct for professional engineers. https://www.bcs.org/media/4411/code-of-conduct.pdf, 2023. Accessed: April 2026.

vLLM Contributors. vllm project repository. https://github.com/vllm-project/vllm, 2024. GitHub repository.

Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for transformer-based generative models. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, pages 521–538. USENIX Association, 2022. URL https://www.usenix.org/conference/osdi22/presentation/yu.