Load by Design: Workload-Shaped Checkpoint Loading for Large Language Model Inference

Botir Khaltaev
Independent, United Kingdom
btrghstk@gmail.com

Abstract

Cold-start latency in large language model (LLM) serving is often treated as a storage-throughput problem: buy faster NVMe and load checkpoints faster. This paper shows that this view is incomplete. Checkpoint layout determines the I/O request geometry exposed to the operating system, backend policy determines how those requests are submitted, and integration layers such as Python and vLLM can amplify or hide backend-level effects. The right loader is therefore workload-shaped, not universal.

We present Tensora, an open measurement framework for checkpoint loading, and use it to evaluate synchronous POSIX I/O, Tokio-based async I/O, Linux `io_uring`, and an adaptive default policy. The study spans two layouts—contiguous SafeTensors shards and a partitioned ServerlessLLM-style layout—and three measurement layers: Rust backend attribution, Python bindings, and a custom vLLM integration. The Rust matrix covers six public checkpoints from 0.6B to 32B parameters. Cross-environment validation on H100, H200, and A100 systems tests whether ordering regimes survive changes in hardware, storage topology, and kernel policy.

Four findings emerge. First, SafeTensors has a scale-dependent crossover: synchronous loading is fastest at small scales, while `io_uring` wins on the H100 host for 8B–32B multi-shard checkpoints. Second, the partitioned ServerlessLLM layout changes the problem: synchronous loading is never the fastest backend in the H100 matrix, and the relevant trade-off shifts to async, `io_uring`, and default read planning. Third, integration changes rankings but does not eliminate loader impact. Python tensor materialisation adds substantial overhead, and in vLLM cold `load_only` runs Tensora's adaptive SafeTensors path reduces initialisation time by 37–48% relative to the native loader for 8B and 32B models. Fourth, regime structure is more portable than backend identity: small SafeTensors models favour sync across all three environments, but the mid-to-large-scale winner depends on whether `io_uring` is available and on storage topology.

The contribution is an evidence-backed policy result: deploy a format-aware adaptive loader with explicit overrides, selected from measured features such as bytes, shard count, partition count, backend availability, and integration layer. Do not choose an I/O API name as a universal default.

***K*eywords** large language models · checkpoint loading · I/O policy · io_uring · cold start · vLLM · reproducibility

1 Introduction

An LLM serving request cannot produce a token until model weights have been made resident. For a 7B model in float16, the checkpoint payload is on the order of 14 GB; for larger models it quickly becomes tens of gigabytes. Serverless and multi-tenant deployments routinely evict idle models to recover GPU memory, so the next request may pay a cold load before inference can start. That load is user-visible and operationally expensive.

The common shorthand is to call this a storage-bandwidth problem. That shorthand hides the mechanism that matters to a loader. A checkpoint format defines the byte ranges that must be read. A backend defines how those ranges are submitted, batched, completed, and copied into buffers. A serving stack then adds Python object construction, tensor materialisation, CUDA initialisation, memory-pool setup, and scheduler overhead. A device with high peak throughput can still deliver poor load time if the request geometry and submission policy are mismatched.

This paper studies the following thesis:

> **LLM checkpoint loading should be selected from workload shape.** Backend policy depends jointly on checkpoint layout, model scale, backend availability, and integration layer. A fixed global rule—always sync, always async, or always `io_uring`—is the wrong abstraction.

Two layouts expose why. SafeTensors checkpoints present mostly contiguous whole-file or multi-shard eager loads. ServerlessLLM-style checkpoints present partitioned range loads with a central index and many grouped reads. These are not small variations of one workload; they induce different I/O regimes. Scale then moves the crossover within a layout: a policy that is effective for a small single-shard load may leave throughput unused on a large multi-shard checkpoint.

We built Tensora[1], an open checkpoint-loading framework and evaluation harness, to measure this effect under controlled and integrated conditions. Tensora exposes synchronous POSIX I/O, Tokio-based async loading, Linux `io_uring`, and a feature-driven `default` policy under common format abstractions. We evaluate six public checkpoints from 0.6B to 32B parameters, two layouts, and three layers: Rust backend attribution, Python bindings, and a custom vLLM integration. The primary environment is an H100 host with local NVMe and working `io_uring`; H200 and A100 container environments with network-attached storage and blocked `io_uring` test capability-gated fallback behaviour.

The data support a policy claim rather than a universal-loader claim. On H100 SafeTensors, explicit sync wins at the small end of the ladder, while explicit `io_uring` wins from Qwen3-8B through Qwen3-32B. On H100 ServerlessLLM, sync is never the fastest backend; async, `io_uring`, and default read planning dominate different rows. Python and vLLM change rankings because they measure different constructs, but loader choice still matters: in vLLM cold `load_only` runs, Tensora's adaptive SafeTensors path reduces initialisation by 37–48% relative to vLLM's native loader for the 8B and 32B checkpoints. Cross-environment results preserve the strongest structural fact—small SafeTensors models favour sync—while showing that the large-scale winner depends on kernel and storage context.

Contributions.

1. A workload-shape characterisation of LLM checkpoint loading: format-level request geometry, not peak bandwidth alone, determines which backend policy is plausible.
2. A controlled Rust-layer measurement matrix comparing sync, async, `io_uring`, and adaptive default across SafeTensors and ServerlessLLM-style layouts.
3. Cross-layer evidence showing how Python bindings and vLLM integration alter, but do not erase, backend effects.
4. Cross-environment validation across H100, H200, and A100 systems showing that ordering regimes are more portable than the identity of the winning backend.
5. A practical policy recommendation: use a format-aware adaptive selector with explicit overrides, capability checks, and measured workload features instead of a universal I/O backend.

Roadmap. Section 2 positions the gap. Section 3 describes Tensora and the adaptive policy. Section 4 defines methodology. Section 5 presents results across three layers. Section 6 interprets findings, Section 7 bounds claims, and Section 8 concludes.

2 Related Work

This section positions our thesis—that LLM checkpoint loading depends on workload shape and should be selected adaptively—against four bodies of work: serving systems, serverless cold-start systems, checkpoint

[1]Tensora is this paper's checkpoint-loading framework (Rust crate `tensora`, Python package `tensora_py`). It is unrelated to Google's TensorStore library for array storage.

formats and loading infrastructure, and operating-system I/O policy. The common gap is that prior work either optimises after weights are resident, avoids loads through caching or placement, or improves the kernel/storage substrate. Tensora measures the application-level load mechanics themselves.

2.1 LLM Serving Systems

Production LLM serving engines primarily optimise throughput and latency once weights are resident. vLLM introduces PagedAttention for efficient KV-cache management [1]. Orca pioneers iteration-level batching [2]. Sarathi-Serve refines the prefill–decode scheduling trade-off [3]. DeepSpeed Inference applies model parallelism and kernel fusion [4]. These systems make the serving pipeline efficient after model initialisation; checkpoint loading is usually treated as a fixed prelude to the serving path.

Our vLLM experiments use a serving engine precisely to test whether backend effects survive that larger context. They do not replace serving-system research; they isolate a cost that serving systems usually inherit from the loader.

2.2 Serverless Cold Start

Serverless LLM serving makes cold start a first-order concern: idle models are evicted and must be reloaded on demand. ServerlessLLM [5] proposes multi-tier loading with GPU/NVMe/DRAM placement and live migration. HydraServe [6] minimises cold-start latency in public clouds through pipeline parallelism and placement optimisation. Snapshot-based warm-up approaches [7] pre-materialise process or model state to bypass part of startup.

These systems optimise *whether*, *where*, or *when* loading occurs. Tensora asks a complementary question: when a load is unavoidable, which userspace submission policy moves the required bytes fastest for this checkpoint geometry? That question remains relevant even in tiered-storage systems because each tier still exposes an I/O workload to the loader.

2.3 Checkpoint Formats and Loading Systems

Checkpoint formats determine what byte ranges a loader must read and whether those reads are contiguous, sharded, partitioned, or fragmented. SafeTensors [8] provides safe tensor serialisation with contiguous whole-file or multi-shard layouts. GGUF optimises for memory-mapped CPU inference and quantised deployments. ServerlessLLM's partitioned layout [9] distributes tensors across partition files with a central index, enabling partial loading and tiered placement.

Recent systems work optimises the loading path from different layers. BlitzScale [10] uses GPU-side model caching over NVLink and RDMA to share already-loaded models between inference instances. Most relevant to our work, PPC/MAIO [11] optimises model loading through a programmable kernel page cache, adapting prefetching and eviction to LLM I/O patterns. MAIO operates below the system-call interface; Tensora operates above it, comparing userspace I/O backends on identical logical read plans. These approaches are complementary: kernel cache policy determines how pages are served, while userspace backend policy determines how application reads are issued.

A separate line of characterization work studies LLM checkpoint/restore I/O patterns on parallel file systems [12], covering buffered versus direct I/O paths and liburing performance. That work focuses on training checkpointing in HPC environments; our focus is inference checkpoint loading in single-node serving.

The gap we address is therefore format-conditioned backend selection: given a fixed checkpoint layout, which userspace backend and planning strategy should execute the read plan, and how does the answer change with model scale and integration layer?

2.4 I/O Backend Policy

The Linux I/O interface landscape spans blocking POSIX calls, asynchronous runtimes built above them, POSIX AIO, and `io_uring`. Recent database research provides the closest methodological template for our work: Jasny et al. [13] show that `io_uring` benefits depend on workload shape and queue depth. Bulk sequential reads may already amortise syscall overhead, while many-small or high-concurrency workloads benefit from batched submission.

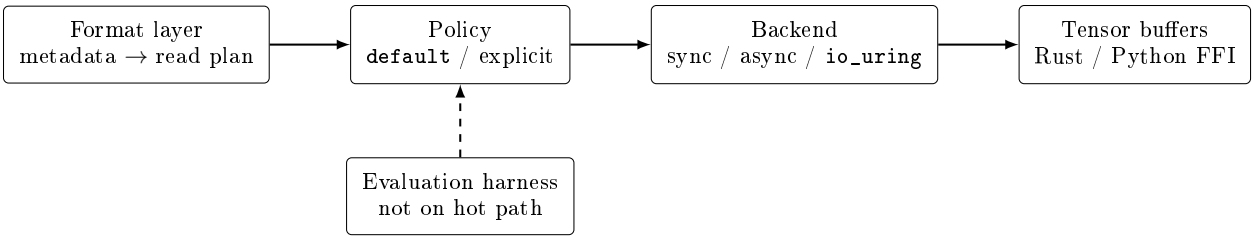

Figure 1: Tensora architecture. The format layer produces a normalised read plan; policy selects a backend and planning parameters; backends execute I/O; evaluation harnesses attach only at the boundary.

LLM checkpoint loading spans both regimes. Contiguous SafeTensors shards resemble bulk sequential access with shard-level fanout. Partitioned ServerlessLLM layouts resemble grouped range workloads. Our work translates the workload-shape lesson to LLM loading and extends it with cross-layer evidence (Rust, Python, vLLM), capability-gated fallback when `io_uring` is blocked, and an adaptive policy grounded in workload features rather than API names.

2.5 Positioning

We do not introduce a new checkpoint format, optimise steady-state serving throughput, or solve cluster scheduling. We measure load mechanics across four backend policies and two layouts, for six public checkpoints, at three integration layers. The contribution is operational: evidence that backend choice depends jointly on layout, scale, and availability, plus a format-aware adaptive policy with explicit overrides. This is the missing piece between format design (what bytes exist), serving engines (how loaded weights are used), and kernel-level storage optimisation (how pages are served).

3 Tensora Design

This section describes *Tensora*, the experimental framework used for the measurement study. The framework is intentionally narrow: it makes checkpoint-loading workload shape observable, runs multiple backends over comparable read plans, and carries the same policy into Python and vLLM integration tests.

3.1 Design Goals

Tensora was designed around four requirements.

1. **Comparable workloads.** Multiple formats and backends must share a common loading abstraction so backend changes do not silently change the logical work.
2. **Separation from measurement.** Cold-cache rituals, timing, and result export belong to the harness, not to the loader.
3. **Explicit concurrency.** Queue depth, worker fanout, chunking, and range grouping are first-class design variables.
4. **Cross-layer reach.** The same core must be callable from Rust benchmarks, Python bindings, and a serving-engine integration.

These requirements led to a layered architecture that separates format parsing, policy selection, I/O execution, and evaluation.

3.2 System Architecture

Tensora has three layers. The **format layer** parses checkpoint metadata and produces a request-local view of the workload: which files and byte ranges must be read to materialise tensors. The **policy/backend layer** chooses a backend and executes that read plan using backend-specific planning such as file chunking, range coalescing, and in-flight depth control. The **evaluation layer** contains the Rust profiler, Python benchmarks, and vLLM integration benchmarks. It observes the boundary; it is not on the hot path.

Figure 1 summarises the control path. The concrete Rust path for the default SafeTensors load is `Model::load(path)` → `choose_load_backend(stats)` → one of `load_sync`, `load_async`, or `load_io_uring`. ServerlessLLM follows the same pattern after compiling an index-derived partition plan.

3.3 Common Backend Interface

All eager-loading backends implement the same conceptual surface:

- `load(path)` for a full file,
- `load_batch(paths)` for multiple whole-file loads,
- `load_range(path, offset, len)` for one byte range, and
- `load_range_batch(requests)` for many ranges.

This interface is deliberately narrow. It covers whole-file SafeTensors loading and range-oriented ServerlessLLM loading without exposing format-specific shortcuts to one backend. The benchmark comparison therefore remains about submission and planning strategy.

3.4 Workload Features

The format layer extracts workload features used both for analysis and for the adaptive policy:

- **Total bytes**: aggregate checkpoint size across shards or partitions.
- **Shard / partition count**: number of discrete files or partition files.
- **Request count**: number of byte-range reads needed to materialise tensors.
- **Contiguity**: whether the load is dominated by whole-file sequential reads or grouped ranges.
- **Format family**: SafeTensors or ServerlessLLM-style partitioning.

SafeTensors checkpoints are characterised by low request count and high contiguity: tensors live in whole files or contiguous shards. ServerlessLLM-style checkpoints have higher partition and range structure: tensors are reconstructed from indexed partition files. These differences are the root cause of the backend-ordering reversals observed in Section 5.

3.5 Adaptive Loading Policy

The central design decision is that `default` is an adaptive heuristic, not an alias for one backend. The Rust matrix (Tables 5 and 6) shows why: no single explicit backend is best across both layouts and all scales.

The `default` policy operates in three stages. First, the format layer computes workload statistics: total bytes, shard or partition count, and coarse contiguity. Second, `choose_load_backend(stats)` maps those features to a backend while consulting `backend_capabilities()` so unavailable interfaces, such as `io_uring` under container security policy, are excluded. Third, backend-specific planning applies chunking, coalescing, and in-flight limits before execution.

This means `default` is not merely a selector among explicit columns. It combines backend choice with execution-time planning. That is intentional: a production loader should both choose a suitable backend and execute it with an appropriate plan. The oracle analysis in Section 5.4 reports where this integrated default beats explicit paths and where it misses.

For `SafeTensors`, the heuristic maps shard count and total bytes to a regime decision: synchronous execution for smaller eager loads, and `io_uring` for larger multi-shard loads when available. If `io_uring` is blocked, large candidates fall back to async. For `ServerlessLLM`, the feature set is partition and request structure; the heuristic does not route to synchronous I/O because the Rust matrix shows sync never wins on this layout.

The selector is feature-driven, not a universal byte threshold. Exact constants are implementation details tied to the measured revision; the paper claims regimes, not magic numbers. Figure 2 summarises control flow. Table 7 quantifies boundedness: `default` is strong in stable regimes and weakest in medium-scale rows where additional signals would be needed.

```
features:  bytes, shards/partitions, contiguity
                          ↓
       capability-gated backend selection
                          ↓
         backend-specific read planning
                          ↓
         execute and materialise tensors
```

Figure 2: Adaptive policy control flow. The selector extracts layout and scale features, gates unavailable backends, then plans and executes reads with backend-specific chunking or range grouping.

Table 1: Major design decisions and rationale.

Decision	Alternatives considered	Rationale
Adaptive `default`	Fixed global backend	Evidence shows crossovers with scale and layout
`io_uring` implementation	Minimal parallelism / naive batching	Persistent rings, batching, and planning are needed to sustain queue depth
Cold-cache control	In-process eviction hints	External `sync`/`drop_caches` separates harness from loader
vLLM integration	Async loader bridge	vLLM's loader API at the Python boundary is synchronous; benchmarks follow that interface
Python Tokio	New runtime per call	Shared runtime reduced binding overhead

3.6 Backend Designs

The synchronous backend uses blocking POSIX reads. It is not a deliberately weak serial baseline: large files can be split into chunks and loaded with host-side thread parallelism. This gives sync a practical implementation and makes its small-model wins meaningful.

The async backend is built on Tokio. Its design is strongest for grouped and range-heavy workloads rather than large whole-file eager loads. It groups work per file, uses dynamic batching, and limits unnecessary task fanout. This explains why async is weak for H100 SafeTensors but often strong for ServerlessLLM and for `io_uring`-restricted environments.

The `io_uring` backend uses persistent rings, batched submissions, lazy initialisation, dynamic ring sizing, and workload-driven planning for chunk size and in-flight depth. Its benefit depends on sustaining enough queued work; the paper does not claim that the API name alone guarantees speed.

Table 2: Evaluated `io_uring` backend structure.

Aspect	Design
Ring ownership	Persistent reader/writer rings with lazy initialisation
Submission	Batched SQEs and explicit in-flight management
Planning	Dynamic ring sizing, chunk sizing, coalescing, and in-flight limits
Fallback	Capability checks exclude `io_uring` when ring creation fails

The mmap-backed path is retained as a useful lazy-access mode, especially for Python and vLLM experiments where a caller may not need to materialise every tensor immediately. It is not a main eager-loading solution to cold start: mmap optimises selective access, while this paper studies aggressive eager loading.

3.7 Checkpoint Formats

Tensora supports two checkpoint layouts. `SafeTensors` stores tensors in a whole-file or multi-shard layout. Once the header is parsed, the loader sees a largely contiguous eager-loading workload. The `ServerlessLLM`

layout distributes tensors across partition files with a central index, creating a fragmented range-oriented workload. Supporting both formats exposes two I/O regimes under one framework.

3.8 Python Bindings and vLLM Integration

Tensora exposes Python bindings implemented with PyO3. The public API intentionally stays small: a `default` eager path, explicit `sync` and `async` eager paths, and lazy `open_*` mmap-backed handles. On Linux, `io_uring` remains behind `default` in the public Python API; benchmark-only explicit variants expose additional backend behaviour for the vLLM layer.

One binding decision was especially important: the async Python path uses a shared Tokio runtime rather than constructing a fresh runtime per call. This reduced binding overhead and improved the public default path in Table 13.

The vLLM integration uses a custom model-loader path. vLLM is a useful evaluation target because it separates model initialisation, time-to-first-token, and steady-state decode. The key constraint is that vLLM's loader interface is synchronous, so the benchmark integration follows that interface rather than inventing an async serving path that the engine does not expose.

Cold-cache control is implemented in the experiment harness, not inside the loader. The Linux shell runs `sync` and `drop_caches` before benchmark invocations. This keeps the loader responsible for loading checkpoints and the harness responsible for creating cold and warm conditions.

4 Experimental Setup

This section describes the experimental methodology used to evaluate the checkpoint loading backends. We make explicit what each measurement layer captures and what it does not capture, so that readers can correctly interpret which claims are supported by which evidence.

4.1 What is measured and why the separation matters

The evaluation spans three distinct measurement layers, each answering a different question:

1. **Rust-layer microbenchmark.** This measures pure I/O and format deserialisation: reading checkpoint bytes from storage through the selected backend and assembling tensors in memory. It intentionally excludes Python interpreter overhead, tensor materialisation in other runtimes, and GPU allocation inside a serving engine. The purpose is backend attribution - isolating which submission policy (sync, async, io_uring, default) produces the best raw loading performance for a given checkpoint layout.
2. **Python-layer binding benchmark.** This measures what a Python client of the Tensora library observes: the full pipeline including FFI transitions, tensor construction, and any binding-specific overhead. It answers whether backend differences survive into the actual API that users would employ. It does not isolate I/O policy from the cost of crossing into Python.
3. **vLLM-layer integration benchmark.** This measures checkpoint loading inside a real serving engine: CUDA context setup, memory pool allocation, runtime configuration, model initialisation, and time-to-first-token. It is the right metric for deployment impact but conflates I/O policy with engine bring-up. It does not answer what pure disk-to-tensor performance is.

These are not interchangeable. A claim that "io_uring is faster" at the Rust layer does not automatically hold at the vLLM layer, because the vLLM layer includes fixed-cost engine initialisation that dominates loader choice for small checkpoints. The limitations section (Section 7) elaborates on this construct validity issue.

4.2 Research Questions

The evaluation is organised around five research questions, each tied to a specific measurement layer:

- **RQ1: Backend attribution.** How much does backend choice affect checkpoint loading time for real LLM checkpoints?

- **RQ2: Layout dependence.** How does checkpoint layout change the relative ordering of backends?
- **RQ3: Public API overhead.** How much overhead is introduced by the Python binding layer, and how much can binding-specific optimisation remove?
- **RQ4: Serving impact.** To what extent do backend gains survive integration into vLLM?
- **RQ5: Operational policy.** What lessons emerge when backend-specific microbenchmarks become a practical adaptive loading system?

RQ1 and RQ2 are answered primarily by Rust-layer matrices. RQ3 is answered by Python-layer benchmarks. RQ4 is answered by vLLM-layer integration. RQ5 draws on all three and is bounded by the limitations in Section 7.

4.3 Experimental Platform

Measurements were collected on three Linux environments to test whether backend ordering regimes transfer across hardware and deployment contexts.

Table 3: Experimental platforms. The H100 host uses local NVMe with full `io_uring` support. The H200 and A100 hosts are network-storage containers where `io_uring` is blocked by a kernel security policy (`EPERM`), exposing the adaptive selector's fallback behaviour.

	H100	**H200**	**A100**
GPU	NVIDIA H100	NVIDIA H200	NVIDIA A100 80GB
CPU	x86_64 multi-core	Intel Xeon 8568Y+	AMD EPYC 7543
RAM	∼180 GB	∼2 TB	∼1 TB
Storage	Local NVMe	Network-attached	Network-attached
`io_uring`	Available	Blocked (EPERM)	Blocked (EPERM)
Kernel	6.x (Ubuntu 24.04)	6.8.0-83 (Ubuntu 22.04)	6.5.0-45 (Ubuntu 22.04)

The H100 host is the primary measurement site: a dedicated Linux server with local NVMe, used as an isolated experiment host with minimal concurrent tenant load. The H200 and A100 environments are containerised GPU instances with network-attached storage. In both containers, the kernel blocks `io_uring` ring creation with `EPERM` (verified via `profile capabilities`). This constraint is not artificial—it reflects a common production deployment pattern where container security policies restrict kernel interfaces. The two restricted environments test whether the adaptive selector correctly falls back from `io_uring` to the next-best backend and whether the resulting ranking still outperforms a fixed policy.

The software stack for final runs was Ubuntu 24.04 (H100) or 22.04 (H200, A100), Python 3.12, and Rust 1.94 (`rustc 1.94.0`, stable). Kernel version was recorded beside exports (`uname -r`). The optional vLLM dependency is pinned in `bindings/python/uv.lock` (resolved `vllm` 0.11.0); serving exports should record `pip show vllm` or `uv pip show vllm` next to tables. Appendix A lists other environment fields to capture.

The H100 host uses local NVMe behind a standard Linux journaling filesystem with default metadata behaviour. The H200 and A100 containers use network-attached storage, which changes the I/O latency profile. Cold-cache runs for the Rust matrix use buffered I/O and external cache drops (procedure below), not O_DIRECT, unless a specific backend documents otherwise. On restricted containers where `drop_caches` is unavailable, the `-evict-page-cache` flag provides in-process page-cache eviction via `posix_fadvise`.

GPU-backed vLLM runs are constrained by vLLM memory budgeting (KV cache, CUDA graphs, buffers), not only raw VRAM. I/O experiments are Linux-only because the io_uring backend is Linux-specific.

4.4 External baselines

The Rust matrices compare Tensora's own backends on *identical* logical read plans; they isolate submission policy rather than pitting Tensora against unrelated third-party loaders.

The vLLM evaluation includes a production-style baseline: the `native` loader uses vLLM's stock checkpoint path for the same model configuration and cache mode. Loaders prefixed `ts_` are Tensora-backed integrations (`ts_safetensors_*`, `ts_serverlessllm_*`) through the same harness so that differences isolate the weight-loading path. Table 14 summarises cold `load_only` times where `native` is available.

At the Python layer, benchmarks target the Tensora public API (explicit backends and `default`). A conventional Hugging Face workflow often uses framework-native tensor construction on top of `safetensors`; that stack differs from our Rust-first path, so we report Tensora binding stages (Table 13) rather than claiming a single universal Python baseline number.

4.5 Backends Under Test

The system exposes four eager-loading policies at the Rust layer:

- **sync**: blocking POSIX reads with dynamic chunking and parallelism.
- **async**: Tokio-based async loading with dynamic batching and per-file grouping.
- **io_uring**: a low-level Linux `io-uring` backend with worker-owned rings, dynamic ring sizing, and dynamic chunk planning (Section 3).
- **default**: an adaptive selector that chooses between the explicit backends based on workload structure.

Rust-layer timings for `io_uring` refer to that implementation (Section 3), not to the syscall family in the abstract.

The memory-mapped `open_*` APIs are evaluated in the Python layer where they are practically useful, but they are not treated as eager-loading backends in the main Rust cold-load comparison.

Two checkpoint layouts were evaluated:

- **SafeTensors**: whole-file or multi-shard eager loading.
- **ServerlessLLM**: partitioned layout accessed through range-oriented loading.

The model ladder consists of six open-access checkpoints:

1. `Qwen/Qwen3-0.6B`
2. `HuggingFaceTB/SmolLM3-3B`
3. `Qwen/Qwen3-4B`
4. `Qwen/Qwen3-8B`
5. `Qwen/Qwen3-14B`
6. `Qwen/Qwen3-32B`

This ladder was chosen to cover small, medium, and large checkpoints while also including one cross-family model (SmolLM3-3B) to reduce the risk of overfitting conclusions to a single vendor family.

The ladder is also what makes the SafeTensors narrative empirically meaningful. A single-point benchmark at one model size cannot justify adaptive policy; it can only justify a constant backend choice. By spanning orders of magnitude in total bytes and shard structure, the experiments expose where synchronous loading stops being the best explicit option and where io_uring becomes the best explicit option in Table 5. That progression is the evidential core behind the adaptive policy claim.

Cold-cache behaviour is the primary condition for the Rust and Python eager-loading measurements. Cold-cache runs use a manual Linux-side procedure before each benchmark invocation (no in-loader eviction):

1. run `sync`
2. run `echo 3 > /proc/sys/vm/drop_caches`
3. invoke the benchmark binary or Python benchmark

Warm-cache conditions skip the cache-drop step and reuse the existing page-cache state.

Rust-layer measurements were generated with a standalone profiling harness. Two categories of Rust experiments were used. First, a five-repetition cold-cache anchor set was collected for Qwen3-0.6B, Qwen3-8B, and Qwen3-32B for both SafeTensors and ServerlessLLM, across sync, async, io_uring, and default. These repeated measurements were used to identify stable workload regimes and to set default-selector thresholds.

Second, a full cold-cache matrix was run across the six-model ladder, both formats, and the four eager backend policies.

Python benchmarks were run against the public binding surface rather than against every internal Rust backend directly. This reflects the actual intended Python API: adaptive default policy, explicit synchronous eager load, explicit Tokio eager load, and mmap-backed lazy access path. The main Python benchmark model was Qwen/Qwen3-8B.

The Python layer was benchmarked across stages that separate pure backend effects from binding-layer tuning and selector behaviour (Table 13).

vLLM benchmarks were run through a custom benchmark loader integration. Unlike the Python eager-loading benchmarks, the vLLM evaluation measures the loading backends inside a wider serving-engine context.

Three benchmark kinds were collected:

- **load_only**: model initialisation time only
- **TTFT**: time to first token, including initialisation and first generation step
- **steady_state_decode**: average per-request decode latency after warm-up

Both warm and cold cache conditions were collected. The full vLLM matrix covers all six model checkpoints and all available benchmark loaders.

Three rules were used to keep the evaluation interpretable.

1. Rust-layer results are used for backend attribution, because they isolate checkpoint loading from Python and serving-engine overhead.
2. Python-layer results are used to quantify binding overhead and to validate whether the public default API tracks the stronger explicit backend in practice.
3. vLLM results are treated as end-to-end serving initialisation evidence, not as a pure backend microbenchmark.

Backend gains do not propagate unchanged across layers.

4.6 Operational definitions and internal validity

What "load time" means at each layer. At the Rust layer, the profiling harness measures the elapsed time to read checkpoint bytes from storage through the selected backend and assemble tensors in memory according to the format's semantics. It intentionally isolates that work from Python object construction and from GPU allocation inside a serving engine. At the Python layer, reported times include interpreter overhead, tensor materialisation, and any FFI transitions exposed by the binding API; they answer whether a real client of the library would observe backend differences, not what a pure Rust microbenchmark would print. At the vLLM layer, `load_only` and `TTFT` include engine initialisation steps beyond tensor bytes on disk: CUDA context setup, memory pools, and runtime configuration. Steady-state decode measures post-warm-up per-request latency and is dominated by compute and scheduling; it is included to show where loader choice stops mattering, not to rank I/O APIs for decode.

Cold cache versus warm cache. A **cold-cache** run is one where the checkpoint bytes are not resident in the Linux page cache at the start of the measurement. The procedure used here is `sync` followed by `echo 3 > /proc/sys/vm/drop_caches`; that sequence drops clean cache pages system-wide before invoking the benchmark. That is stronger than evicting a single file with advisory hints and matches common practice in storage evaluation. It does not guarantee an empty GPU or an idle disk firmware queue; it guarantees that the loader cannot satisfy reads entirely from DRAM backed by a prior read of the same file. A **warm-cache** run skips the drop and reuses pages left by earlier reads on the host, modelling a second load without a full cache flush between runs.

Why eviction is out-of-process. In-process `posix_fadvise` and similar hints are advisory and filesystem-dependent; embedding them in the loader would blur the line between harness and code under test. External cache drops keep the loader binary unchanged: every cold iteration pays the same visible storage path, and the methodology is easy to replicate from a shell script.

Variance and anchor repetitions. NVMe and kernel subsystems exhibit run-to-run jitter. Single cold numbers are therefore *representative* rather than statistically exhaustive. The five-repetition anchor set exists to confirm that observed crossovers are not one-off artefacts for a handful of ladder points (small, large, and mid-scale anchors across formats). The full matrices use one cold iteration per cell for readability. Interpreting a 5% gap as meaningful without repeats is risky; interpreting a consistent ordering across many cells and formats with anchors showing stable bands is much less so.

Comparability across backends. Each backend receives the same logical tensor read plan for a given model and format: the same sequence of byte ranges must be satisfied. Differences in wall time come from submission strategy, parallelism, and buffer management, not from silently skipping work. When `default` loses to an explicit backend on a row, that is an empirical fact about the selector and planner under the same harness constraints, not a different problem instance.

Statistical reporting. Full matrices use representative cold-cache times; anchor cells use five repetitions to check variance. Appendix B reports median wall-clock time and the min–max band for selected anchor configurations (SafeTensors and ServerlessLLM). Main tables omit per-cell standard deviations for readability, so small gaps in one cell should be read cautiously, while consistent ordering across many cells carries more weight.

Cross-environment scope. The primary Rust and vLLM matrices come from the H100 host with local NVMe and full `io_uring` support. The H200 and A100 environments provide cross-hardware validation under a different constraint: network-attached storage with `io_uring` blocked. We do not claim that absolute timings transfer across machines. The paper claims that **ordering regimes** transfer more plausibly; the restricted environments test that claim directly by removing `io_uring` from the available backend set and observing whether the adaptive selector produces reasonable fallback orderings. Section 7 elaborates.

Linux-only scope. The io_uring backend is Linux-specific. Results on other operating systems would require a different I/O backend and are outside scope.

4.7 Claim boundaries and falsification criteria

The central claim is that backend choice depends jointly on layout and scale, so a fixed global backend is wrong in principle. The evidence for this is the crossover pattern: sync wins at small SafeTensors scale, io_uring wins at large SafeTensors scale, and sync never wins for partitioned ServerlessLLM.

This claim would be falsified if:

- Explicit synchronous loading remained fastest for the largest SafeTensors checkpoints in the ladder (opposite of observed)
- Explicit io_uring never won at large SafeTensors scale (opposite of observed)
- Synchronous loading was fastest across the entire ServerlessLLM matrix (opposite of observed)

The claim would be supported (but boundary location refined) if replication on different hardware showed the same ordering pattern but at different crossover points in the model ladder. The H200 and A100 environments partially test this: they remove `io_uring` from the available set and show that the adaptive selector falls back to appropriate alternatives. Section 7 states these falsification conditions explicitly.

The adaptive `default` policy claim would be falsified if the selector systematically lost to a single explicit backend across the full matrix. The evidence shows it wins on some rows and loses on others (medium-scale SafeTensors, notably SmolLM3-3B and Qwen3-4B). This is documented in Section 7.

vLLM results would not falsify the Rust-layer I/O claims because they conflate loader choice with engine bring-up. They instead test whether backend differences survive into deployment, which is a different question. The limitations section makes this construct distinction explicit.

4.8 Reading the matrices and reproducing ordering

Read Rust tables alongside the results in Section 5. SafeTensors compares sync, async, io_uring, and `default` on the same logical workload; `default` shows when integrated planning beats or misses explicit backends. The ladder orders increasing bytes and shard complexity in this harness only. SmolLM3-3B and Qwen3-4B sit in a medium band where `default` still trails explicit sync (Section 7). On ServerlessLLM, treat sync as a baseline, not a target.

Cold-cache runs use `sync` and `drop_caches` so numbers reflect storage, not warm DRAM. vLLM warm-cache cases are labelled separately.

Replication should not chase identical milliseconds across machines; it should reproduce **ordering**: SafeTensors sync→io_uring crossover among explicit backends, and ServerlessLLM with sync off the front of Table 6. Ship revision hashes, model IDs, partition conversion settings, and cold-cache commands. The adaptive policy is conditional on the measured crossover structure; if ordering differs on other hardware, that is a finding about boundary stability, not a refutation of the methodology.

5 Results

The evaluation proceeds in three layers: Rust backend attribution (Section 5.3), Python binding overhead (Section 5.6), and vLLM integration (Section 5.7). Tables present core evidence; figures visualise regime crossovers. Section 5.4 compares the adaptive default against the best per-row explicit backend to bound the selector's effectiveness.

5.1 Reading the Results

The figures and tables are organised by control flow. First, the Rust layer isolates backend attribution for the same logical read plans. Second, the oracle analysis asks how close the adaptive `default` comes to the best explicit backend. Third, cross-environment results test capability-gated fallback. Finally, Python and vLLM results show how the same loader choices behave once integration overhead is included.

Table 4: Regime map for the strongest backend by layout and environment. Colour encodes the winning explicit or integrated policy in the corresponding matrix. The visual point is the regime shift: no colour dominates every row.

Layer / layout	0.6B	4B	8B	32B
H100 SafeTensors	`sync`	`sync`	`io_uring`	`io_uring`
H100 ServerlessLLM	`io_uring`	`async`	`default`	`default`
H200 SafeTensors	`sync`	`sync`	`async`	`async`
H200 ServerlessLLM	`async`	`async`	`async`	`default`
A100 SafeTensors	`sync`	`async`	`async`	`default`
A100 ServerlessLLM	`default`	`default`	`default`	`async`
vLLM SafeTensors	`sync`	`default`	`default`	`io_uring`

Table 4 is the paper's main visual summary. SafeTensors begins blue (sync) and shifts with scale; ServerlessLLM is not sync-led; restricted environments replace `io_uring` with async or default; vLLM preserves the broad small-versus-large pattern while adding engine overhead.

5.2 Headline findings

Three patterns emerge from the Rust matrices. **First, SafeTensors exhibits a backend crossover.** Synchronous loading is the fastest explicit backend for model sizes up to 4B parameters; io_uring is the fastest explicit backend from 8B parameters upward. No single fixed backend tracks both regimes.

Second, the partitioned ServerlessLLM layout reverses the ranking. Synchronous loading is never the fastest explicit backend for any model size in this ladder. The relevant trade-off is between async loading (strong on smaller and mid-sized range loads) and io_uring (competitive at larger partition counts).

Third, the adaptive default is useful but bounded. It matches or beats the best explicit backend in 5 of 12 H100 Rust cells (Table 7) and stays near parity in several others; it trails significantly on medium-scale rows where the feature mapping is imprecise. This boundedness is informative: it shows where explicit overrides add value.

Across Python and vLLM, integration overhead does not erase backend effects—it reorders them. This validates the operational conclusion: deploy a format-aware adaptive default with explicit overrides rather than a fixed single backend.

Fourth, the ordering regime transfers across environments but the winning backend does not. Cross-environment validation on H200 and A100 containers (where `io_uring` is blocked by kernel security

policy) confirms that the structural pattern—small models favour sync, mid-to-large models typically favour a different backend—holds on H100 and H200, while A100 shows a non-monotonic pattern where async dominates mid-range but sync recovers at the largest scale (Section 5.5). This makes the adaptive selector more valuable, not less: a fixed policy would require manual tuning per deployment site.

5.3 Rust Backend Results

The Rust harness is the authoritative backend-attribution layer because it measures checkpoint loading directly, without Python tensor conversion or serving-engine initialisation costs. Unless stated otherwise, each matrix cell is one representative cold-cache run; Appendix B reports five-repetition medians and min–max bands for selected anchors.

5.3.1 SafeTensors Results

Table 5 summarises the explicit backend comparison for SafeTensors.

Table 5: Rust cold-cache SafeTensors results (ms). The crossover from sync-favored small models to io_uring-favored large multi-shard models is the evidential core behind the adaptive policy.

Model	Sync	Async	Default	io_uring
Qwen3-0.6B	555.28	1416.19	333.59	1159.24
SmolLM3-3B	1331.79	5802.05	3722.35	3812.35
Qwen3-4B	1820.85	5039.21	3083.18	2967.42
Qwen3-8B	3880.40	7074.40	3844.35	3172.37
Qwen3-14B	7030.25	13489.00	6138.70	5977.07
Qwen3-32B	15454.45	28763.46	12119.50	12089.05

Two regimes are visible. Sync dominates at the small end: 555 ms for Qwen3-0.6B versus 1159 ms for io_uring. The ranking flips at Qwen3-8B: io_uring at 3172 ms versus sync at 3880 ms, a 1.22× advantage. At the largest scale (Qwen3-32B), io_uring is 1.28× faster than sync. The crossover is not a single threshold; it is a progression visible across the entire ladder.

Async is consistently the weakest backend for SafeTensors. Whole-file and multi-shard loads benefit from straightforward bulk parallelism rather than scheduler-mediated async orchestration, which adds overhead without corresponding I/O benefits on this workload shape.

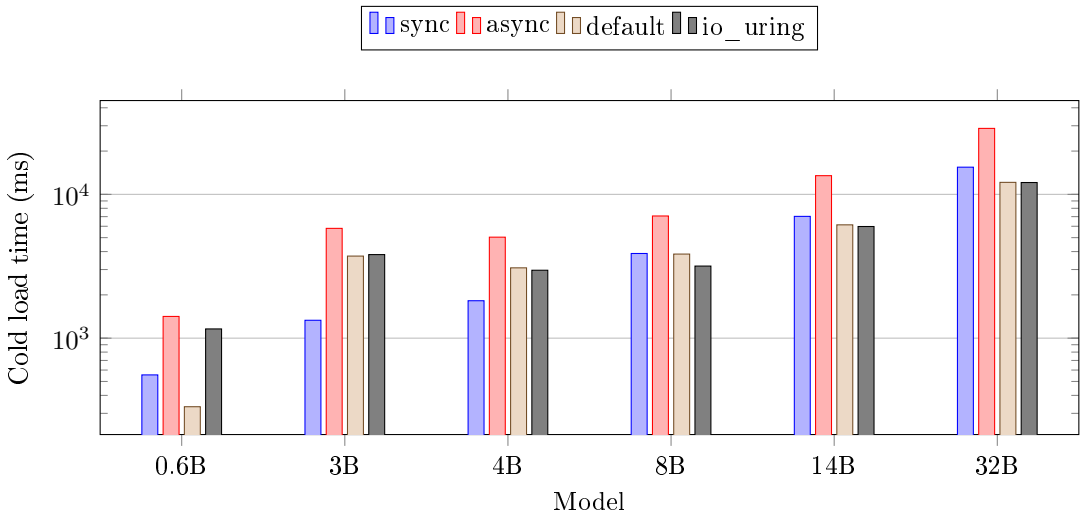

Figure 3: SafeTensors backend crossover (log scale). Sync dominates the left (small models); io_uring dominates the right (large multi-shard models). Default trails on medium-scale rows where the heuristic's feature mapping is imprecise.

Figure 3 visualises the progression. `default` adds planning (chunking, coalescing) that helps on Qwen3-0.6B (333 ms, beating all explicit backends) but misses the medium-scale crossover region on SmolLM3-3B and Qwen3-4B.

5.3.2 ServerlessLLM Results

The ServerlessLLM layout produces a different ordering, shown in Table 6.

Table 6: Rust cold-cache ServerlessLLM results (ms). Synchronous loading never wins; the ranking shifts from async at smaller partition scales toward io_uring at larger sizes.

Model	Sync	Async	Default	io_uring
Qwen3-0.6B	1645.06	1471.34	1432.25	1387.57
SmolLM3-3B	3941.88	2391.24	3182.28	3091.49
Qwen3-4B	5137.72	3504.25	3551.82	3727.14
Qwen3-8B	11048.40	9463.37	7589.47	9043.51
Qwen3-14B	18661.25	16139.08	13422.94	13257.23
Qwen3-32B	39822.09	29180.89	27659.78	29779.86

Sync is the slowest explicit backend on every ServerlessLLM row, between 1.1× (Qwen3-0.6B) and 1.4× (Qwen3-32B) behind the strongest alternative. The best explicit backend varies with scale: async dominates at smaller partition counts (SmolLM3-3B, Qwen3-4B), while io_uring becomes competitive at larger scales (Qwen3-14B, Qwen3-32B).

`default` outperforms all explicit backends on Qwen3-8B (7589 ms vs 9044 ms for the next best, io_uring) and Qwen3-32B (27660 ms vs 29181 ms for async) because its read planning—coalescing nearby ranges, dynamic chunk sizing—benefits partitioned workloads more than any single explicit submission policy. It trails explicit async on SmolLM3-3B (3182 ms vs 2391 ms, a 1.33× gap), the same medium-scale fragility seen in SafeTensors.

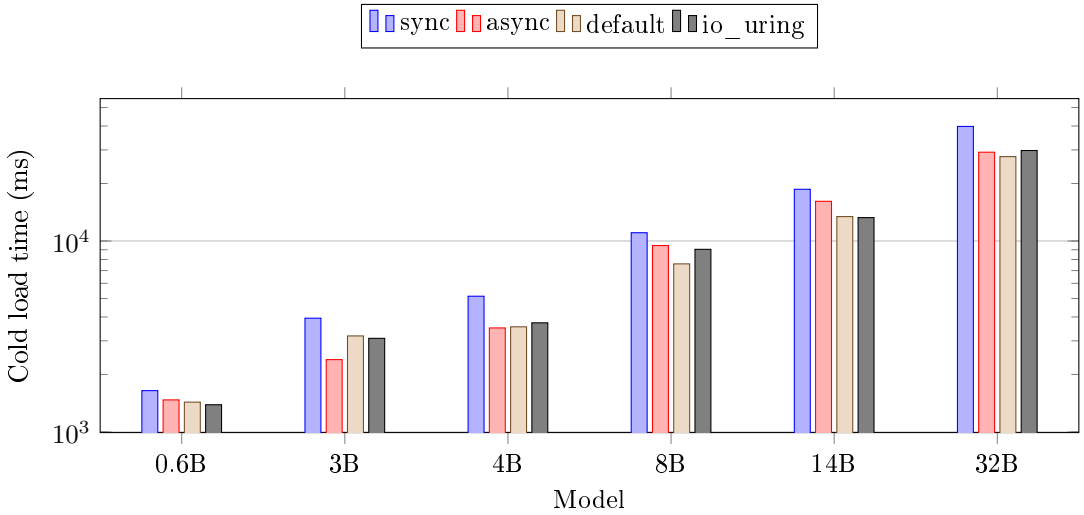

Figure 4: ServerlessLLM backend ordering (log scale). Sync is never the winner; the layout shifts the ranking toward async and io_uring. Default benefits from read planning on large partitioned loads.

Figure 4 shows the ranking shift. The figure confirms that the ServerlessLLM workload shape—many small range reads across partition files—rewards different backend policies than the SafeTensors contiguous regime.

5.4 Oracle analysis: how well does the adaptive default track the best backend?

The adaptive default is not expected to match the best per-row explicit backend in every cell—it is a heuristic, not an oracle. Table 7 quantifies the gap.

Table 7: Default versus per-row best explicit backend. Cells show `default` / `best_explicit` ratio. Ratios < 1.0 mean default wins; ratios > 1.0 mean default trails.

Model	SafeTensors	ServerlessLLM
Qwen3-0.6B	0.60	1.03
SmolLM3-3B	2.79	1.33
Qwen3-4B	1.69	1.01
Qwen3-8B	1.21	0.80
Qwen3-14B	1.03	1.01
Qwen3-32B	1.00	0.93

Default matches or beats the best explicit backend in 5 of 12 cells (ratios ≤ 1.0). For SafeTensors, default correctly routes the smallest model (Qwen3-0.6B) to a fast path and tracks the large-scale io_uring regime within 3% at Qwen3-14B and near parity at Qwen3-32B. The medium-scale misses—SmolLM3-3B (2.79×), Qwen3-4B (1.69×), and Qwen3-8B (1.21×)—are the region where the crossover from sync to io_uring is most sensitive to per-checkpoint shard structure that the current feature set does not fully resolve.

For ServerlessLLM, default matches or beats the best explicit backend in 3 of 6 cells. The single notable miss is SmolLM3-3B (1.33×), where async's grouped range execution fits the partition structure better than default's chosen path. The remaining ratios are near-parity (≤ 1.03).

Interpretation. The misses are not hidden; they are the point at which the adaptive policy becomes an engineering policy rather than an oracle. They identify where the feature mapping is fragile and where explicit overrides (for example, force sync for SafeTensors SmolLM3-3B; force async for ServerlessLLM SmolLM3-3B) provide clear value. A deployable selector should therefore expose overrides instead of pretending every boundary can be learned from one static threshold.

5.5 Cross-Environment Validation

The H100 results above were collected with all four backends available, including `io_uring`. Two additional environments—an H200 container and an A100 container, both with network-attached storage—block `io_uring` at the kernel level (`EPERM`). This constraint tests two questions: (1) whether the adaptive selector degrades gracefully when its strongest large-scale backend is unavailable, and (2) whether backend ordering regimes transfer across hardware.

5.5.1 H200 Results (io_uring blocked)

Table 8 shows the H200 SafeTensors matrix. With `io_uring` blocked, async replaces it as the fastest backend at large scale, while sync remains dominant for small models.

Table 8: H200 cold-cache SafeTensors results (ms). `io_uring` is blocked (`EPERM`). Sync dominates at small scale; async takes over at large scale in the absence of `io_uring`.

Model	Sync	Async	Default
Qwen3-0.6B	604.34	852.17	649.23
SmolLM3-3B	2013.66	2938.30	2130.51
Qwen3-4B	909.83	2478.10	941.24
Qwen3-8B	3924.69	2977.59	3213.01
Qwen3-14B	25872.03	4767.67	4906.13
Qwen3-32B	36205.91	10623.83	12356.23

The H200 SafeTensors crossover preserves the structural pattern from H100: sync wins at small scale (0.6B–4B), while a different backend dominates at large scale (8B+). On H100 that backend is `io_uring`; on H200 it is async. The selector adapts: `default` tracks within 3–8% of the best available backend for 0.6B–14B, though it trails async by 16% at 32B.

Table 9 shows the H200 ServerlessLLM matrix.

Table 9: H200 cold-cache ServerlessLLM results (ms). Async dominates most rows; `default` beats both explicit backends at 32B.

Model	Sync	Async	Default
Qwen3-0.6B	1565.27	1491.44	1586.84
SmolLM3-3B	3945.68	1947.39	1987.07
Qwen3-4B	5100.25	2489.27	2809.79
Qwen3-8B	10788.88	6045.47	6098.06
Qwen3-14B	17765.28	14450.40	15453.47
Qwen3-32B	39061.57	27367.74	24259.15

The ServerlessLLM ordering on H200 matches H100 structurally: sync is never the fastest, and async dominates most rows. `default` outperforms both explicit backends at 32B (24259 ms vs 27368 ms for async), consistent with the H100 finding that read planning benefits partitioned loads at large scale.

5.5.2 A100 Results (io_uring blocked, reduced ladder)

The A100 environment runs a four-model reduced ladder (0.6B, 4B, 8B, 32B) with the same `EPERM` constraint. Table 10 and Table 11 summarise the results.

Table 10: A100 cold-cache SafeTensors results (ms). Network-attached storage with AMD EPYC CPU. Sync wins at 0.6B; async dominates at 4B–8B; sync recovers at 32B.

Model	Sync	Async	Default
Qwen3-0.6B	724.26	1094.43	965.95
Qwen3-4B	7598.41	3532.11	7810.53
Qwen3-8B	8412.36	4235.21	11896.57
Qwen3-32B	46923.42	60807.34	43816.68

Table 11: A100 cold-cache ServerlessLLM results (ms). `default` wins at 0.6B, 4B, and 8B; async dominates at 32B.

Model	Sync	Async	Default
Qwen3-0.6B	2183.33	2250.91	1554.97
Qwen3-4B	7551.79	4486.75	3523.25
Qwen3-8B	14707.01	11528.26	11431.61
Qwen3-32B	56025.98	40472.54	44854.60

The A100 SafeTensors results show a distinctive pattern: sync wins at 0.6B, async dominates the mid-range (4B–8B), but sync recovers at 32B. This contrasts with both H100 (where `io_uring` dominates large scale) and H200 (where async dominates large scale). The `default` selector on A100 performs well at 0.6B and 32B but misses the mid-range async regime, producing its largest gap at 8B (2.81× behind async). For ServerlessLLM, `default` wins at 0.6B and 4B and tracks within 1% of async at 8B.

5.5.3 Cross-Environment Summary

Table 12 summarises the best explicit backend per row across all three environments.

Three findings from the cross-environment comparison: (1) Sync dominates small SafeTensors models in all three environments—this is the most portable ordering. (2) At large scale, the strongest backend depends on available kernel interfaces and storage topology: `io_uring` on local NVMe (H100), async on network storage (H200), and sync on a different network-storage configuration (A100 at 32B). (3) The adaptive selector degrades gracefully: when `io_uring` is unavailable, it produces within-regime-competitive results on most rows, confirming that the capability-gated fallback mechanism works as designed.

Table 12: Best explicit backend by environment for SafeTensors. The crossover from sync to another backend persists on H100 and H200; A100 shows a non-monotonic pattern where sync recovers at 32B. Which alternative wins at mid-to-large scale depends on hardware and storage topology.

Model	H100	H200	A100
Qwen3-0.6B	sync	sync	sync
Qwen3-4B	sync	sync	async
Qwen3-8B	io_uring	async	async
Qwen3-32B	io_uring	async	sync

5.6 Python Results

Python benchmarks were collected on Qwen/Qwen3-8B to quantify user-facing overhead and to evaluate whether binding optimisation improved practical performance.

The most important Python result is that binding engineering matters. Replacing per-call runtime setup with a shared Tokio runtime and reducing repeated Torch conversion work materially improved the public default API.

Table 13: Python eager-load times for Qwen3-8B (seconds). Binding tuning and selector redesign both contributed measurable improvements.

Case	Baseline bindings	After binding tuning	After selector tuning
SafeTensors default	15.88	14.55	11.55
ServerlessLLM default	22.45	20.50	18.45

Table 13 shows the sequence. For SafeTensors on Qwen3-8B, the default path moves from 15.88 s to 11.55 s, a 27.3% reduction from the baseline. For ServerlessLLM, the same sequence moves default from 22.45 s to 18.45 s, a 17.8% reduction.

The Python numbers are much larger than the Rust-only numbers (11.55 s vs 3.84 s for SafeTensors default) because eager loading at the binding layer includes substantial tensor materialisation and Python object construction. The key conclusion: Python is not a transparent pass-through layer. Backend improvements matter, but they are filtered through the cost of converting checkpoint contents into Torch tensors and Python dictionaries.

5.7 vLLM Results

The final vLLM matrix completed successfully across all six models, all benchmark loaders, warm and cold cache conditions, and all three benchmark kinds.

Stock vLLM versus Tensora. The `native` loader follows vLLM's default weight-loading path; `ts_*` loaders use Tensora-backed integrations. On cold `load_only`, Tensora's default SafeTensors path reduces initialisation time relative to `native` across all measured scales (Table 14).

Table 14: Cold vLLM load_only: native vs. Tensora (ms). The gap is an end-to-end effect (CUDA context, memory pools, and loader policy), not a pure Rust microbenchmark.

Model	`native`	`ts_safetensors_default`	Δ
Qwen3-0.6B	23659	20578	−13.0%
Qwen3-8B	56639	35556	−37.2%
Qwen3-32B	169971	88200	−48.1%

The improvement grows with model size: 13% for the smallest model, 37% at 8B, and 48% at 32B. This is consistent with the Rust SafeTensors finding (io_uring advantage grows at larger scales) filtered through vLLM's fixed initialization overhead.

For cold load_only across all models, the best loader varies with model size. Table 15 shows the strongest loader in the final cold-load-only configuration.

Table 15: Best cold load_only vLLM loader by model. The pattern follows the Rust SafeTensors regime: sync-style loading at small scales, io_uring at large scales.

Model	Best init time (ms)	Best loader
Qwen3-0.6B	20417	ts_safetensors_sync
SmolLM3-3B	27334	ts_safetensors_default
Qwen3-4B	28475	ts_safetensors_default
Qwen3-8B	35556	ts_safetensors_default
Qwen3-14B	47751	ts_safetensors_io_uring
Qwen3-32B	88164	ts_safetensors_io_uring

The vLLM best-loader pattern mirrors the Rust SafeTensors crossover: smaller models favour sync-style loaders, larger models favour default or io_uring. The match is not exact—vLLM initialization includes fixed costs that dilute the pure-I/O ranking—but the regime structure survives.

For TTFT (Table 16), the ranking is broadly similar to load_only but does not match row-for-row because TTFT includes the first generation step after initialisation.

Table 16: Best cold TTFT vLLM loader by model.

Model	Best TTFT (ms)	Best loader
Qwen3-0.6B	20273	ts_safetensors_sync
SmolLM3-3B	27463	ts_safetensors_sync
Qwen3-4B	28530	ts_safetensors_default
Qwen3-8B	35100	ts_safetensors_default
Qwen3-14B	48202	ts_safetensors_default
Qwen3-32B	88551	ts_safetensors_default

The SafeTensors loader family consistently outperforms the ServerlessLLM family for both metrics. Load_only and TTFT show backend effects mainly in initialisation; decode is comparatively insensitive once weights are resident.

5.8 Layer summary

Four principles emerge. (1) Backend ranking depends on workload shape: SafeTensors favours sync at small scale and `io_uring` at large H100 scale, while ServerlessLLM favours async, `io_uring`, or default planning and disfavours sync. (2) The adaptive default is useful but bounded: it wins or ties several stable-regime rows and misses medium-scale rows that need explicit overrides. (3) Integration layers reorder but do not eliminate backend differences: Python adds tensor conversion overhead and vLLM adds initialisation cost, yet loader choice still changes user-visible startup. (4) Cross-environment validation confirms that sync dominates small SafeTensors models everywhere; the mid-to-large-scale winner varies with kernel interface availability and storage topology, and A100 shows a non-monotonic pattern where sync/default recovers at the largest scale. The selector's capability-gated fallback adapts to these constraints better than a hard-coded `io_uring`-first policy would.

6 Discussion

The central lesson is simple: checkpoint loading is a workload-shape problem. Layout determines request geometry; scale determines how much queue depth and parallelism can be used; deployment context determines which kernel interfaces are available; integration layers determine how much backend improvement reaches the user.

6.1 Why the layouts diverge

SafeTensors and ServerlessLLM-style partitioning expose different I/O problems. SafeTensors is mostly whole-file or multi-shard eager loading. For small checkpoints, a practical sync implementation is hard to

beat: it has low orchestration overhead and enough parallelism for the available work. As checkpoint size and shard count grow, keeping more I/O in flight matters, and the H100 matrix shifts toward `io_uring`.

ServerlessLLM-style partitioning is different. The loader reconstructs tensors from partition files and range-oriented access. This makes grouping, coalescing, and batching central. In this regime the sync-first intuition from SafeTensors fails: sync is never the fastest backend in the H100 ServerlessLLM matrix. The best policy is therefore not "use `io_uring`" or "use sync"; it is "read the format and scale before choosing."

6.2 Why integration changes the ranking

Rust results answer the backend-attribution question. Python and vLLM answer different questions. Python adds interpreter overhead, tensor conversion, and object materialisation. vLLM adds CUDA context setup, memory pools, runtime configuration, and generation setup. These costs can dilute or reorder pure I/O differences.

That reordering is not a weakness in the evaluation; it is the deployment reality. A loader that wins only in a microbenchmark but disappears at the serving boundary is less valuable. Conversely, a loader that saves tens of seconds in vLLM initialisation is operationally meaningful even if the exact Rust ranking changes.

6.3 What the adaptive policy should promise

The adaptive selector should promise three things and no more. First, it should avoid obviously wrong global defaults, such as sync for all ServerlessLLM loads or `io_uring` for every small SafeTensors load. Second, it should degrade gracefully when an interface is unavailable, as in H200 and A100 containers where `io_uring` ring creation is blocked. Third, it should expose explicit overrides for boundary cases.

It should not promise universal optimality. Table 7 shows medium-scale misses, and those misses are useful: they reveal which features are insufficient. A production policy can incorporate more signals over time, but the paper's evidence already rules out a single fixed backend.

6.4 External validity and replication

Absolute milliseconds are the least portable part of the study. Ordering regimes are more portable. The H200 and A100 results support this distinction: small SafeTensors loads favour sync across all environments, while the mid-to-large-scale winner changes with storage topology and `io_uring` availability. A replication should therefore ask whether the same regime changes appear, not whether a different host reproduces every number.

6.5 Design implications

For practitioners, the implication is to log and select on workload features: format family, total bytes, shard or partition count, backend availability, and integration layer. For researchers, the implication is to report loading results as layered measurements. A Rust microbenchmark, a Python API benchmark, and a serving-engine cold-start benchmark are all useful, but they are not interchangeable constructs.

7 Limitations

The results support a regime-based policy claim in a controlled harness. They do not establish a universally optimal loader.

7.1 External validity

The primary Rust and vLLM matrices come from one H100 host with local NVMe and working `io_uring`. H200 and A100 containers provide partial cross-environment validation under a different deployment condition: network-attached storage with `io_uring` blocked by kernel policy. This is useful because it tests fallback behaviour, but it is not the same as repeating the full H100 experiment on multiple local-NVMe hosts.

The structural pattern is stronger than the exact crossover point. Small SafeTensors loads favour sync across all three environments. Larger SafeTensors loads favour `io_uring` on H100, async on H200, and sync/default

at the largest A100 row. The identity of the winner therefore depends on kernel interface availability, storage topology, filesystem behaviour, and host load.

The model ladder covers six public checkpoints from 0.6B to 32B parameters and two layouts. It does not exhaust other formats such as GGUF, quantised checkpoint variants, distributed filesystems, RAID configurations, object stores, or non-Linux operating systems.

7.2 Construct validity

Each measurement layer defines "loading time" differently. Rust timings approximate backend-attributed checkpoint read and tensor assembly. Python timings add FFI, interpreter overhead, tensor materialisation, and Python object construction. vLLM timings include GPU bring-up, memory pools, runtime configuration, model initialisation, and first-token generation depending on the benchmark kind. None of these exactly equals production user-perceived latency, which may add queuing, batching, network RTT, model routing, and admission control.

For this reason, claims do not move freely across layers. A Rust result explains backend mechanics; a vLLM result explains serving-startup impact under that harness.

7.3 Selector boundedness

The adaptive default does not match the best per-row explicit backend across all cells. On H100 SafeTensors it trails explicit sync by 2.79× for SmolLM3-3B and 1.69× for Qwen3-4B, and trails explicit `io_uring` by 1.21× for Qwen3-8B. On H100 ServerlessLLM it trails explicit async by 1.33× for SmolLM3-3B. These rows sit near regime boundaries or partition structures that the current feature set does not fully capture.

The right response is not to hide these misses. The operational recommendation is an adaptive default with explicit overrides. The default automates stable regimes; overrides handle known boundary cases.

7.4 Statistical limits

The main matrices report representative cold-cache runs. Anchor repetitions in Appendix B validate selected regimes, but not every cell has repeated statistics. Small gaps below roughly 10% should be interpreted cautiously unless they appear consistently across rows, layers, or repetitions.

7.5 Scope of claims

The evidence supports an engineering heuristic justified by measurement: backend policy should use format, scale, and capability signals. It does not prove optimal thresholds, a learned selector, or universal superiority of `io_uring`. Future work should add more local-NVMe hosts, filesystem and mount-option sweeps, additional checkpoint formats, direct-I/O variants, object-store paths, and calibrated or learned selectors.

8 Conclusion

LLM checkpoint loading is not governed by storage bandwidth alone. It is governed by workload shape: the layout of bytes on disk, the scale of the checkpoint, the backend interfaces available on the host, and the overhead imposed by the integration layer.

This paper measured that shape with Tensora across two checkpoint layouts, four backend policies, three measurement layers, and three hardware environments. The evidence rejects a fixed global backend. SafeTensors shifts from sync-favoured small loads to `io_uring`-favoured large H100 loads. ServerlessLLM-style partitioning changes the ranking so sync is never fastest in the H100 matrix. Python and vLLM reorder the backend story but preserve deployment relevance: loader choice affects public API time and serving initialisation. Cross-environment validation shows that regimes transfer better than exact winners.

The practical conclusion is therefore not "use `io_uring`" or "use sync." It is: load by design. Select a backend from measured workload features, gate it by runtime capability, and expose explicit overrides for boundary cases. A format-aware adaptive policy is the right abstraction for checkpoint loading because the workload itself changes the answer.

A Reproducibility

This appendix summarises how measurements were produced: what to hold constant, where artefacts live, and which commands reproduce the Rust-layer backend ordering. The goal is traceability and comparable ordering, not byte-identical millisecond reproduction on other hardware.

A.1 Source bundle and public PDF

For public distribution, submit `main.tex`, `template.tex`, `PRIMEarxiv.sty`, the `sections/` tree, `figures/`, `library.bib`, and `00README.json`. Build from `paper/` with:

```
pdflatex -interaction=nonstopmode main.tex
bibtex main
pdflatex -interaction=nonstopmode main.tex
pdflatex -interaction=nonstopmode main.tex
```

If a conference-supplied document class is adopted later, follow that venue's template and submission rules separately.

A.2 Repository and revision

The implementation and harnesses are open source at `https://github.com/botirk38/tensora`. The Rust crate `tensora`, Python package `tensora_py`, and benchmark scripts are distributed under the Apache License, Version 2.0. Commands, default output paths, and a minimal environment block are in the repository root `README.md`. Tabulated timings should correspond to a specific commit hash; record `git rev-parse HEAD` and `git status -sb` next to exported results.

A.3 Rust profiling harness

Build the profiling binary in release mode:

```
cargo build --release --bin profile
```

Invoke the Rust layer with a public Hugging Face model id:

```
cargo run --release --bin profile -- <format> <backend> \
  --model-id <MODEL_ID> --iterations 1
```

where `<format>` is `safetensors` or `serverlessllm`, and `<backend>` is `sync`, `async`, `io-uring`, or `default`. SafeTensors shards are read from the Hugging Face Hub cache. ServerlessLLM-style conversions are stored under the platform cache directory (for example `$XDG_CACHE_HOME/tensora/<slug>/serverlessllm`).

Cold-cache runs use `-iterations 1` with the cache-drop procedure in Section 4. The published H100 cold matrix (`results/h100/rust.tsv`) was produced by looping formats, model ids, and backends with `sync` and `echo 3 > /proc/sys/vm/drop_caches` before each `profile` invocation. On restricted containers where `drop_caches` is unavailable, use the repository-supported cache-eviction flag and record that deviation beside the result.

Five-repetition anchors for selected models (`results/h100/anchor.tsv` and `results/h200/anchor.tsv` where available) used the same binary and cache-control ritual repeated five times per anchor. Main tables use representative single cold-cache rows for readability.

A.4 Python and vLLM harnesses

Python benchmarks call the public binding API under `bindings/python/benchmarks/`. The repository provides `scripts/run\_benchmarks.sh`, which runs pytest-benchmark suites and writes JSON artefacts under `results/benchmarks/` by default. Set `TENSORA_BENCH_MODELS` to one or more Hugging Face model ids, separated by spaces or commas. Set `TENSORA_BENCH_NO_VLLM=1` to skip vLLM when a GPU-backed serving environment is unavailable.

When vLLM benchmarks are enabled, the script serialises vLLM runs to one model at a time. Archived exports used in the paper appear under `results/h100/`, `results/h200/`, and `results/a100/`. The H100

vLLM matrix in `results/h100/vllm.tsv` was produced on a dedicated host by sweeping loaders, model ids, benchmark kinds (`load_only`, `TTFT`, `steady_state_decode`), and cache modes.

A.5 Environment pins

Capture at least `date -u`, `uname -a`, `rustc -version`, `python3 -version`, `nvidia-smi` GPU and driver output, and the active vLLM package version. Kernel and CUDA versions affect `io_uring` behaviour and GPU initialisation. The repository lockfile resolves `vllm` to 0.11.0 when installed through `uv`; still record the active environment's package line with exports so a reader can match wheels and CUDA builds.

A.6 What replication should check

On different machines, prioritise ordering and regime behaviour: SafeTensors crossover with scale; Serverless-LLM without sync at the front of the ranking; Python overhead large enough to matter; vLLM initialisation sensitive to loader choice while decode is not. Matching every millisecond is unnecessary for validating the main findings.

A.7 Hub names and cache control

The main text uses public hub names such as `Qwen/Qwen3-8B`. Scripts derive filesystem slugs from those ids. Mixing ad hoc checkpoint paths can change shard counts and invalidate comparisons. Prefer hub checksum verification when available.

Dropping the page cache requires `/proc/sys/vm/drop_caches`, typically with elevated privileges. The experiments assumed a dedicated host so cache drops do not disturb unrelated tenants. Run `sync` before `drop_caches` so cold reads exercise storage rather than stale clean buffers.

B Cold-cache anchor statistics

Table 17 summarises five cold-cache repetitions for selected anchors; source files include `results/h100/anchor.tsv` and `results/h200/anchor.tsv`. Median wall-clock time is reported with the min–max band across repetitions. These anchors sanity-check variance; the main matrices still use one representative cold run per cell for presentation.

Table 17: Anchor repetitions: median cold load time (ms) with min–max across five runs.

Fixture	Format	Backend	Median	Min–Max
Qwen3-0.6B	SafeTensors	sync	323.35	322.05–402.93
Qwen3-8B	SafeTensors	sync	3700.68	3647.97–3731.11
Qwen3-8B	SafeTensors	io_uring	3250.29	3160.46–3274.00
Qwen3-32B	SafeTensors	sync	15462.39	15326.77–16143.61
Qwen3-32B	SafeTensors	io_uring	12204.11	11736.60–14388.76
Qwen3-0.6B	ServerlessLLM	default	1356.17	1336.94–1404.77
Qwen3-8B	ServerlessLLM	default	8059.26	7586.28–9143.26

C Artefact layout

Each environment directory (`results/h100/`, `results/h200/`, `results/a100/`) contains `rust.tsv` for Rust-layer cold-cache profiles, `vllm.tsv` for vLLM integration benchmarks where available, `anchor.tsv` for repetition anchors where available, `env.txt` for environment metadata where recorded, and `raw/` for pytest-benchmark JSON artefacts.

References

[1] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP 2023)*. ACM, 2023.

[2] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for transformer-based generative models. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, pages 521–538. USENIX Association, 2022.

[3] Amey Agrawal, Nitin Kedia, Ashish Panwar, et al. Taming throughput-latency tradeoff in LLM inference with Sarathi-Serve. In *Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 2024)*. USENIX Association, 2024.

[4] Reza Yazdani Aminabadi et al. DeepSpeed inference: Enabling efficient inference of transformer models at unprecedented scale. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC 2022)*. IEEE/ACM, 2022.

[5] Yao Fu, Leyang Xue, Yeqi Huang, Andrei-Octavian Brabete, Dmitrii Ustiugov, Yuvraj Patel, and Luo Mai. Serverlessllm: Low-latency serverless inference for large language models. In *Proceedings of the 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 2024)*. USENIX Association, 2024.

[6] Anonymous Authors. HydraServe: Minimizing cold start latency for serverless LLM serving in public clouds. arXiv preprint arXiv:2502.15524, 2025.

[7] Paulo Silva and Thiago Emmanuel Pereira. Performance evaluation of snapshot methods to warm the serverless cold start, 2021.

[8] Hugging Face. Safetensors documentation. `https://huggingface.co/docs/safetensors`, 2024. Accessed: April 2026.

[9] ServerlessLLM Contributors. Serverlessllm. `https://github.com/ServerlessLLM/ServerlessLLM`, 2024. GitHub repository.

[10] Zizhao Mo et al. Blitzscale: Efficient large model serving with elastic GPU resource sharing. In *Proceedings of the 2024 USENIX Annual Technical Conference (ATC 2024)*. USENIX Association, 2024.

[11] Yubo Liu, Hongbo Li, Xiaojia Huang, Yongfeng Wang, Hanjun Guo, Hui Chen, Yuxin Ren, and Ning Jia. Accelerating model loading in LLM inference by programmable page cache. In *Proceedings of the 24th USENIX Conference on File and Storage Technologies (FAST 2026)*. USENIX Association, 2026.

[12] Nicolas Pretet et al. Understanding LLM checkpoint/restore I/O strategies and patterns, 2025.

[13] Matthias Jasny, Muhammad El-Hindi, Tobias Ziegler, Viktor Leis, and Carsten Binnig. High-performance DBMSs with io_uring: When and how to use it, 2025.